

# Distributed Scheduling with Apache Mesos in the Cloud

PhillyETE - April, 2015  
Diptanu Gon Choudhury  
@diptanu

NETFLIX

OSS

NETFLIX

DEXTER

# Who am I?

- Distributed Systems/Infrastructure Engineer in the Platform Engineering Group
  - Design and develop resilient highly available services
  - IPC, Service Discovery, Application Lifecycle
- Senior Consultant at ThoughtWorks Europe
- OpenMRS/RapidSMS/ICT4D contributor




# A word about Netflix

## Just the stats

- 16 years
- < 2000 employees
- 50+ million users
- $5 * 10^9$  hours/quarter
- Freedom and Responsibility Culture



NETFLIX



# The Titan Framework

A globally distributed resource scheduler which offers compute resources as a service

NETFLIX






# Guiding Principles

Design for

- Native to the public clouds
- Availability
- Reliability
- Responsiveness
- Continuous Delivery
- Pushing to production faster

NETFLIX



# Guiding Principles

- Being able to sleep at night even when there are partial failures.
- Availability over Consistency at a higher level
- Ability for teams to fit in their domain specific needs





NETFLIX

# Active-Active Architecture



# Current Deployment Pipeline

NETFLIX





# The Base AMI

NETFLIX

Base AMI

java

tomcat

httpd

python

others

Foundation Image

# Need for a Distributed Scheduler

- ASGs are great for web services but for processes whose life cycle are controlled via events we needed something more flexible
- Cluster Management across multiple geographies
- Faster turnaround from development to production



# Need for a Distributed Scheduler


- A runtime for polyglot development
- Tighter Integration with services like Atlas, Scryer etc



# We are not alone in the woods

- Google's Borg and Kubernetes
- Twitter's Aurora
- Soundcloud's Harpoon
- Facebook's tupperware
- Mesosphere's Marathon

NETFLIX





# Why did we write Titan

- We wanted a cloud native distributed scheduler
- Multi Geography from the get-go
- A meta scheduler which can support domain specific scheduling needs
  - Work Flow systems for batch processing workloads
  - Event driven systems
  - Resource Allocators for Samza, Spark, etc

# Why did we write Titan

- Persistent Volumes and Volume Management
- Scaling rules based on metrics published by the kernel
- Levers for SREs to do region failovers and shape traffic globally



# Compute Resources as a service

```
{  
  "name": "rocker",  
  "applicationName": "nf-rocker",  
  "version": "1.06",  
  "location": "dc1:20,us-west-2:dc2:40,dc5:60",  
  "cpus": 4,  
  "memory": 3200,  
  "disk": 40,  
  "ports": 2,  
  "restartOnFailure": true,  
  "numRetries": 10,  
  "restartOnSuccess": false  
}
```

# Things Titan doesn't solve

- Service Discovery
- Distributed Tracing
- Naming Service

NETFLIX





# Building blocks


- A resource allocator
- Packaging and isolation of processes
- Scheduler
- Distribution of artifacts
- Replication across multiple geographies
- AutoScalers

NETFLIX



# Resource Allocator

- Scale to 10s of thousands of servers in a single fault domain
- Does one thing really well
- Ability to define custom resources
- Ability to write flexible schedulers
- Battle tested





# Mesos



Apache  
**MESOS**™

NETFLIX

DEXTER

# How we use Mesos

- Provides discovery of resources
- We have written a scheduler called Fenzo
- An API to launch tasks
- Allows writing executors to control the lifecycle of a task
- A mechanism to send messages





# Packaging and Isolation

- We love Immutable Infrastructure
- Artifacts of applications after every build contains the runtime
- Flexible process isolation using cgroups and namespaces
- Good tooling and distribution mechanism

# Docker



NETFLIX






# Building Containers

- Lots of tutorials around docker helped our engineers to pick the technology very easily
- Developers and build infrastructure uses the Docker cli to create containers.
- The docker-java plugin allows developers to think about their application as a standalone process

# Volume Management

- ZFS on linux for creating volumes
- Allows us to clone, snapshot and move around volumes
- The zfs toolset is very rich
- Hoping for a better libzfs

NETFLIX





# Networking

- In AWS EC2 classic containers use the global network namespace
- Ports are allocated to containers via Mesos
- In AWS VPC, we can allocate an IP address per container via ENIs

NETFLIX

Bill & Ted's  
EXCELLENT  
adventure

office

DEXTER

# Logging

- Logging agent on every host to allows users to stream logs
- Archive logs to S3
- Every container gets a volume for logging



# Monitoring

- We push metrics published by the kernel to Atlas
- The scheduler gets a stream of metrics from every container to make scheduling decisions
- Use the cgroup notification API to alert users when a task is killed






# Scheduler

- Remembers the cluster state
  - Efficient bin-packing
  - Helps with Auto Scaling
  - Allows us to do things like reserve instances for specific type of workloads


NETFLIX



# Auto Scaling

- A must need for running on the cloud
- Two levels of scaling
  - Scaling of underlying resources to match the demands of processes
  - Scaling the applications based on metrics to match SLAs

NETFLIX





# Reactive Auto Scaling

- Titan adjusts the size of the fleet to have enough compute resources to run all the tasks
- Autoscaling Providers are pluggable



NETFLIX



# Predictive Autoscaling

- Historical data to predict the size of clusters of individual applications
- Linear Regression models for predicting near real time cluster sizes

NETFLIX





# Bin Packing for efficient Autoscaling

Node A




16 CPUs

Service A

Service A

Service A

Node B




16 CPUs

Long Running Service

Short Lived Batch Process

Short Lived Batch Process

Node C



16 CPUs


# Bin Packing for efficient Autoscaling

Node A




16 CPUs

Node B



16 CPUs

Node C



16 CPUs




# Mesos Framework

- Master Slave model with leader election for redundancy
- A single Mesos Framework per fault domain
- We currently use Zookeeper but moving to Raft
- Resilient to failures of underlying data store

# Globally Distributed

- Each geography has multiple fault domains
- Single scheduler and API in each fault domain.

NETFLIX







NETFLIX

Spinnaker

Apollo Creed

Dagobah

Meson

Samza

Titan

Mesos



Amazon EC2



Amazon EC2



Amazon EC2



# Future

- More robust scheduling decisions
- Optimize the host OS for running containers
- More monitoring

NETFLIX



# Questions?

NETFLIX

