



Apache Hadoop YARN

Enabling next generation data applications

Arun C. Murthy
Founder & Architect
[@acmurthy](#) ([@hortonworks](#))



Hello!

- **Founder/Architect at Hortonworks Inc.**

- Lead - Map-Reduce/YARN/Tez
- Formerly, Architect Hadoop MapReduce, Yahoo
- Responsible for running Hadoop MapReduce as a service for all of Yahoo (~50k nodes footprint)

- **Apache Hadoop, ASF**

- Frmr. VP, Apache Hadoop, ASF (Chair of Apache Hadoop PMC)
- Long-term Committer/PMC member (full time >6 years)
- Release Manager for hadoop-2.x

Agenda

- Why YARN?
- YARN Architecture and Concepts
- Building applications on YARN
- Next Steps

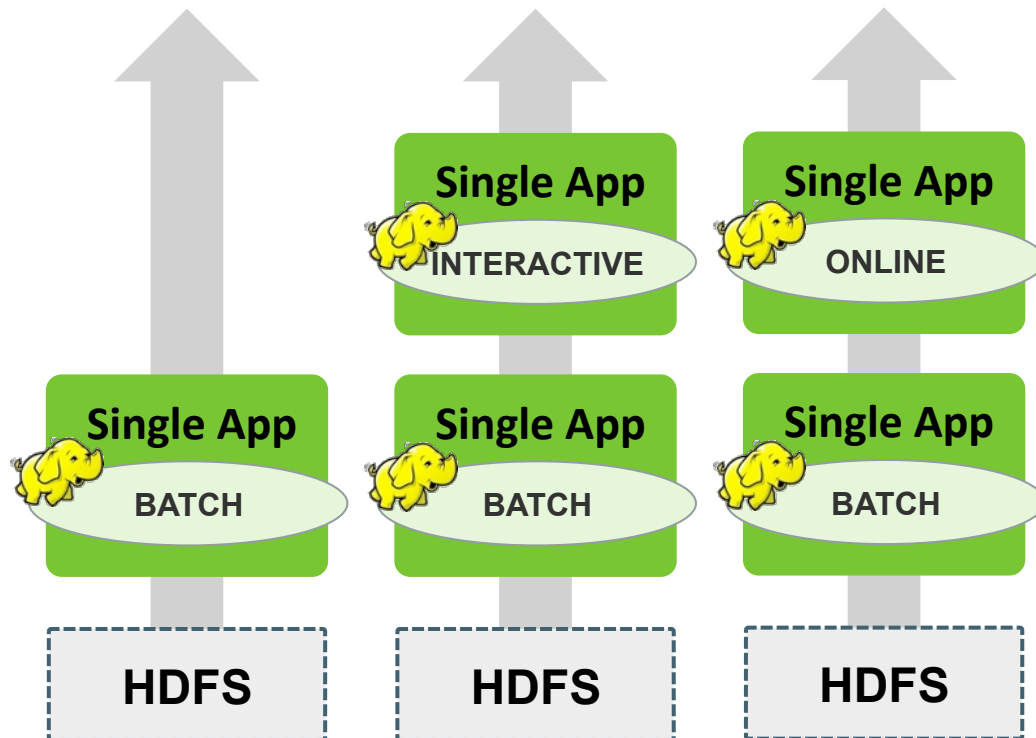
Agenda

- **Why YARN?**
- YARN Architecture and Concepts
- Building applications on YARN
- Next Steps

The 1st Generation of Hadoop: Batch

HADOOP 1.0

Built for Web-Scale Batch Apps



- All other usage patterns must leverage that same infrastructure
- Forces the creation of silos for managing mixed workloads

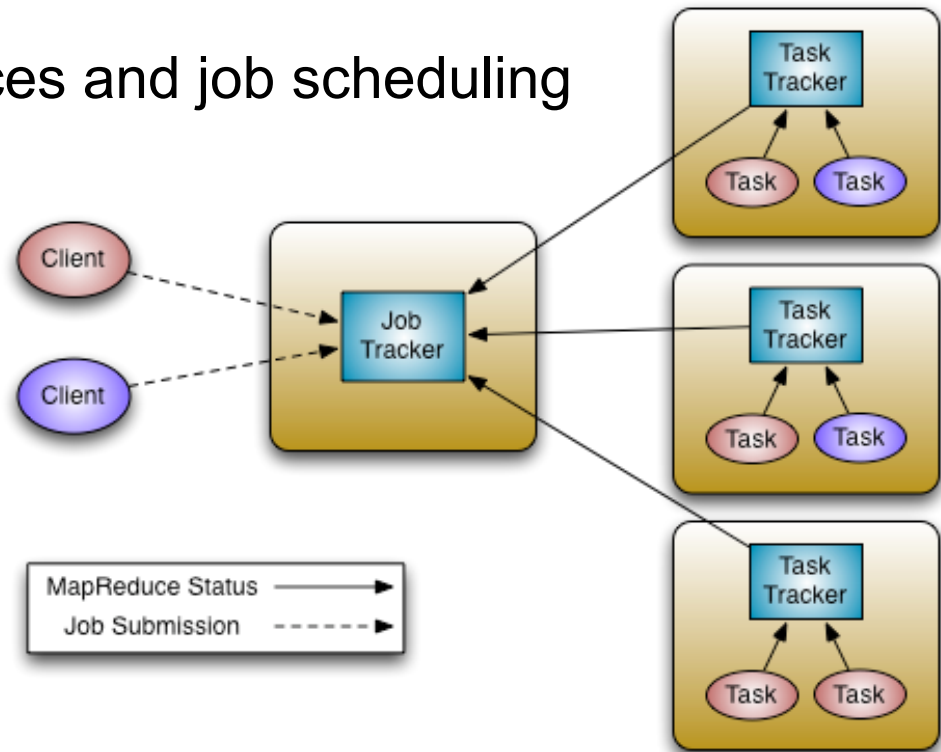
Hadoop MapReduce Classic

- **JobTracker**

- Manages cluster resources and job scheduling

- **TaskTracker**

- Per-node agent
 - Manage tasks



MapReduce Classic: Limitations

- **Scalability**

- Maximum Cluster size – 4,000 nodes
- Maximum concurrent tasks – 40,000
- Coarse synchronization in JobTracker

- **Availability**

- Failure kills all queued and running jobs

- **Hard partition of resources into map and reduce slots**

- Low resource utilization

- **Lacks support for alternate paradigms and services**

- Iterative applications implemented using MapReduce are 10x slower

Our Vision: Hadoop as Next-Gen Platform

Single Use System

Batch Apps

HADOOP 1.0

MapReduce

(cluster resource management
& data processing)

HDFS

(redundant, reliable storage)



Multi Purpose Platform

Batch, Interactive, Online, Streaming, ...

HADOOP 2.0

MapReduce

(data processing)

Others

(data processing)

YARN

(cluster resource management)

HDFS2

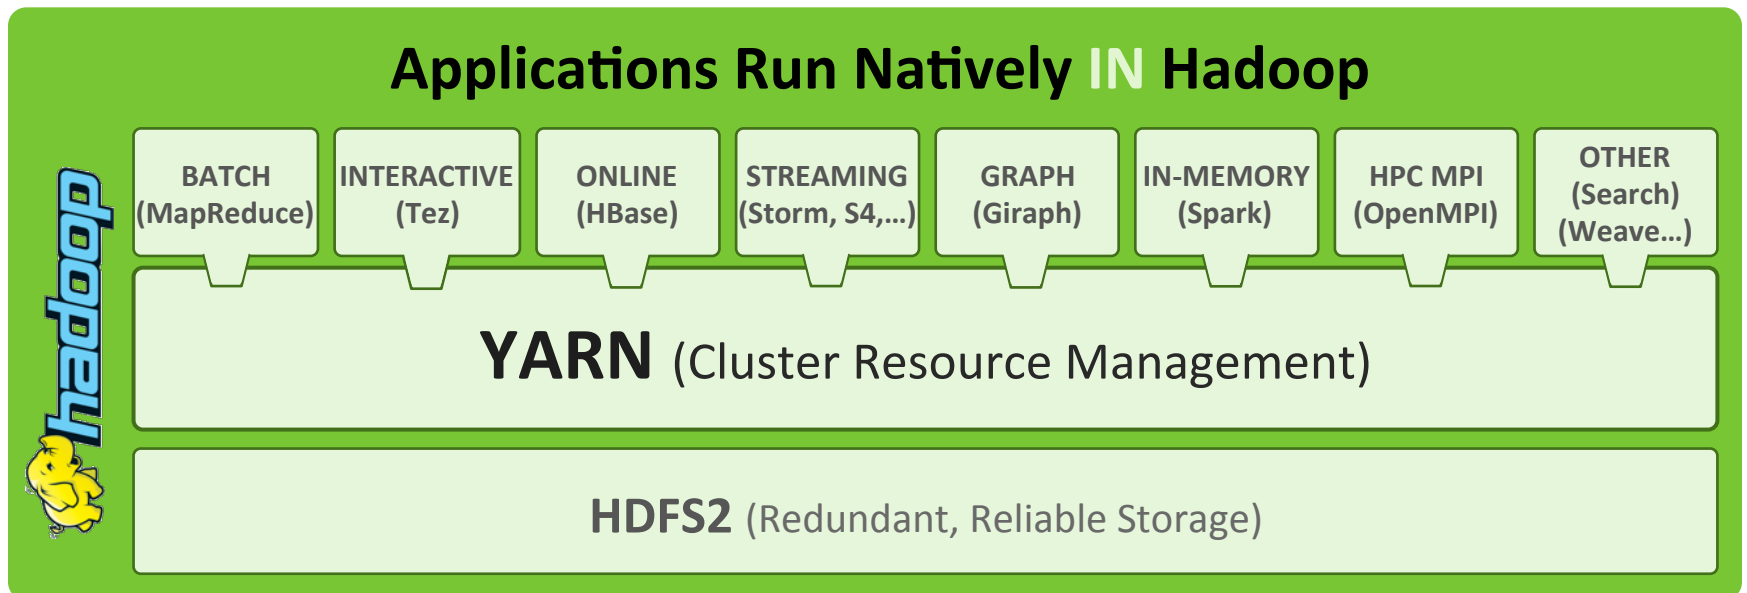
(redundant, reliable storage)

YARN: Taking Hadoop Beyond Batch

Store ALL DATA in one place...

Interact with that data in MULTIPLE WAYS

with Predictable Performance and Quality of Service



5 Key Benefits of YARN

- 1. Scale**
- 2. New Programming Models & Services**
- 3. Improved cluster utilization**
- 4. Agility**
- 5. Beyond Java**

Agenda

- Why YARN
- **YARN Architecture and Concepts**
- Building applications on YARN
- Next Steps

A Brief History of YARN

- **Originally conceived & architected by the team at Yahoo!**
 - Arun Murthy created the original JIRA in 2008 and led the PMC
- **The team at Hortonworks has been working on YARN for 4 years**
- **YARN based architecture running at scale at Yahoo!**
 - Deployed on 35,000 nodes for 6+ months
- **Multitude of YARN applications**

Concepts

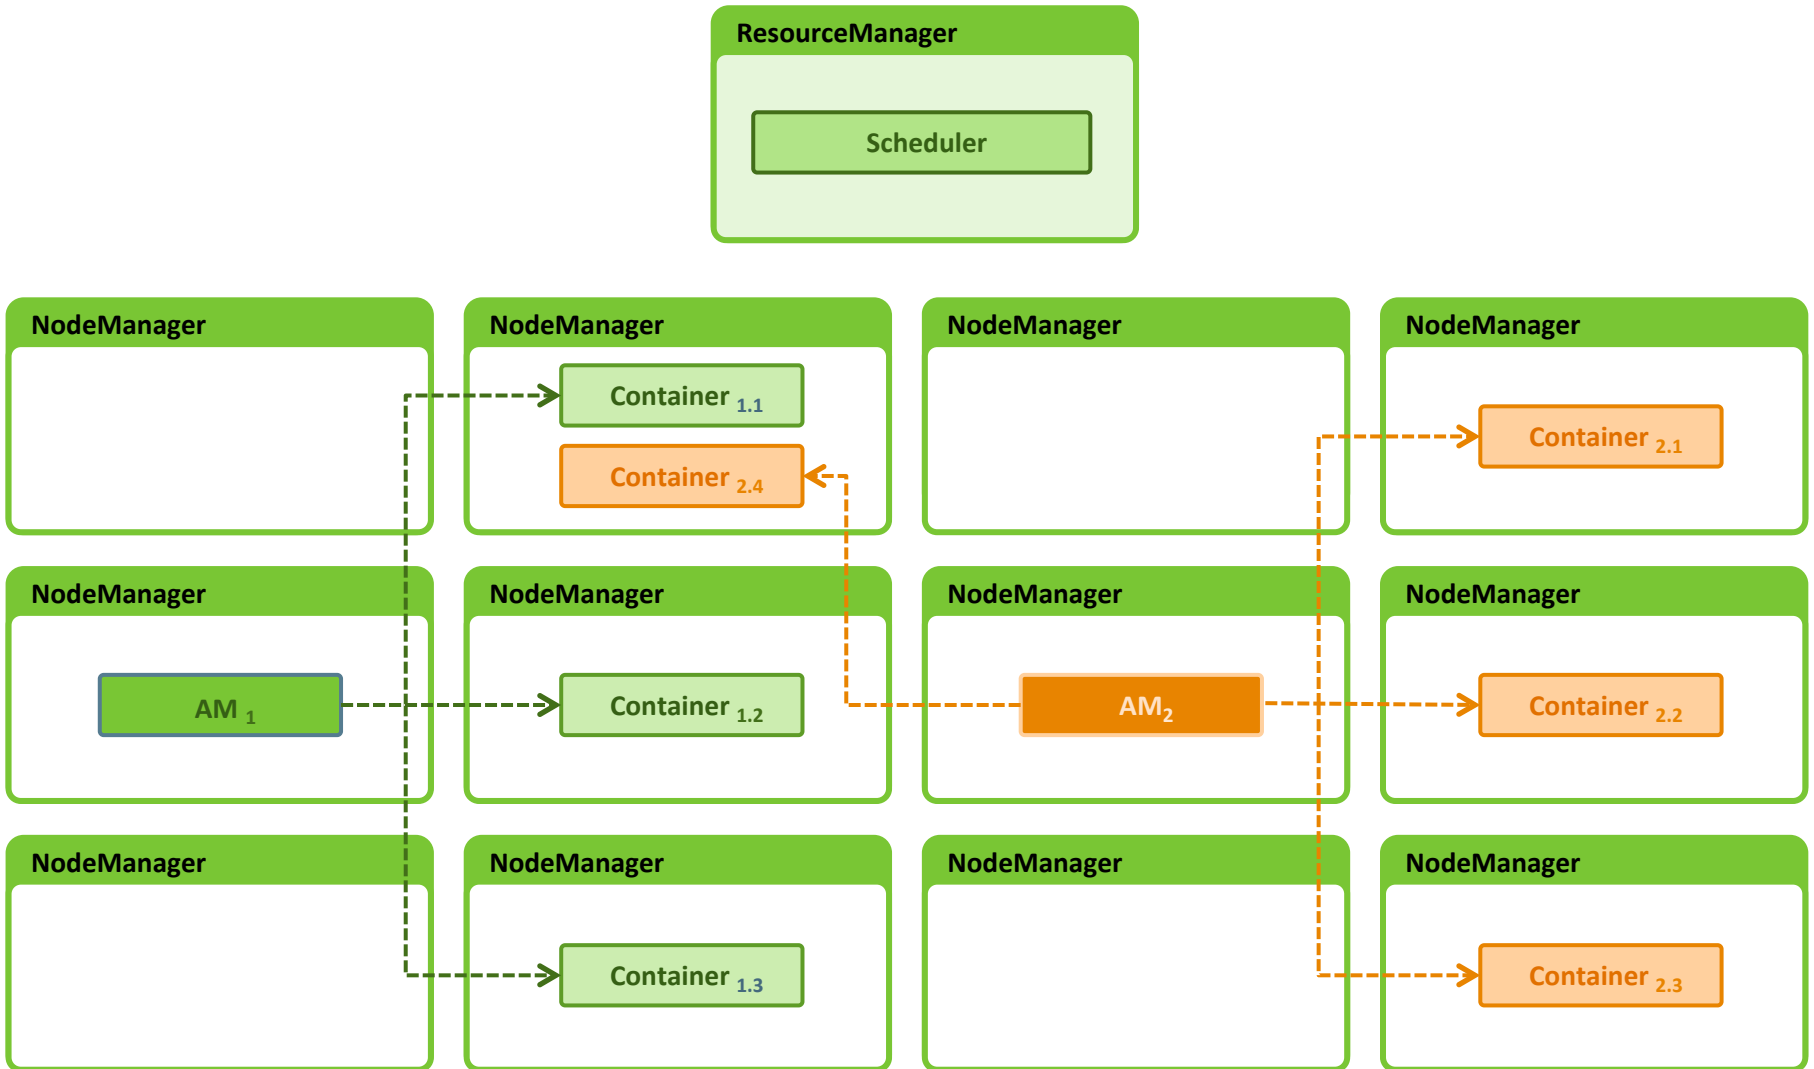
- **Application**

- Application is a job submitted to the framework
- Example – Map Reduce Job

- **Container**

- Basic unit of allocation
- Fine-grained resource allocation across multiple resource types (memory, cpu, disk, network, gpu etc.)
 - container_0 = 2GB, 1CPU
 - container_1 = 1GB, 6 CPU
- Replaces the fixed map/reduce slots

YARN Architecture



Architecture

- **Resource Manager**

- Global resource scheduler
- Hierarchical queues

- **Node Manager**

- Per-machine agent
- Manages the life-cycle of container
- Container resource monitoring

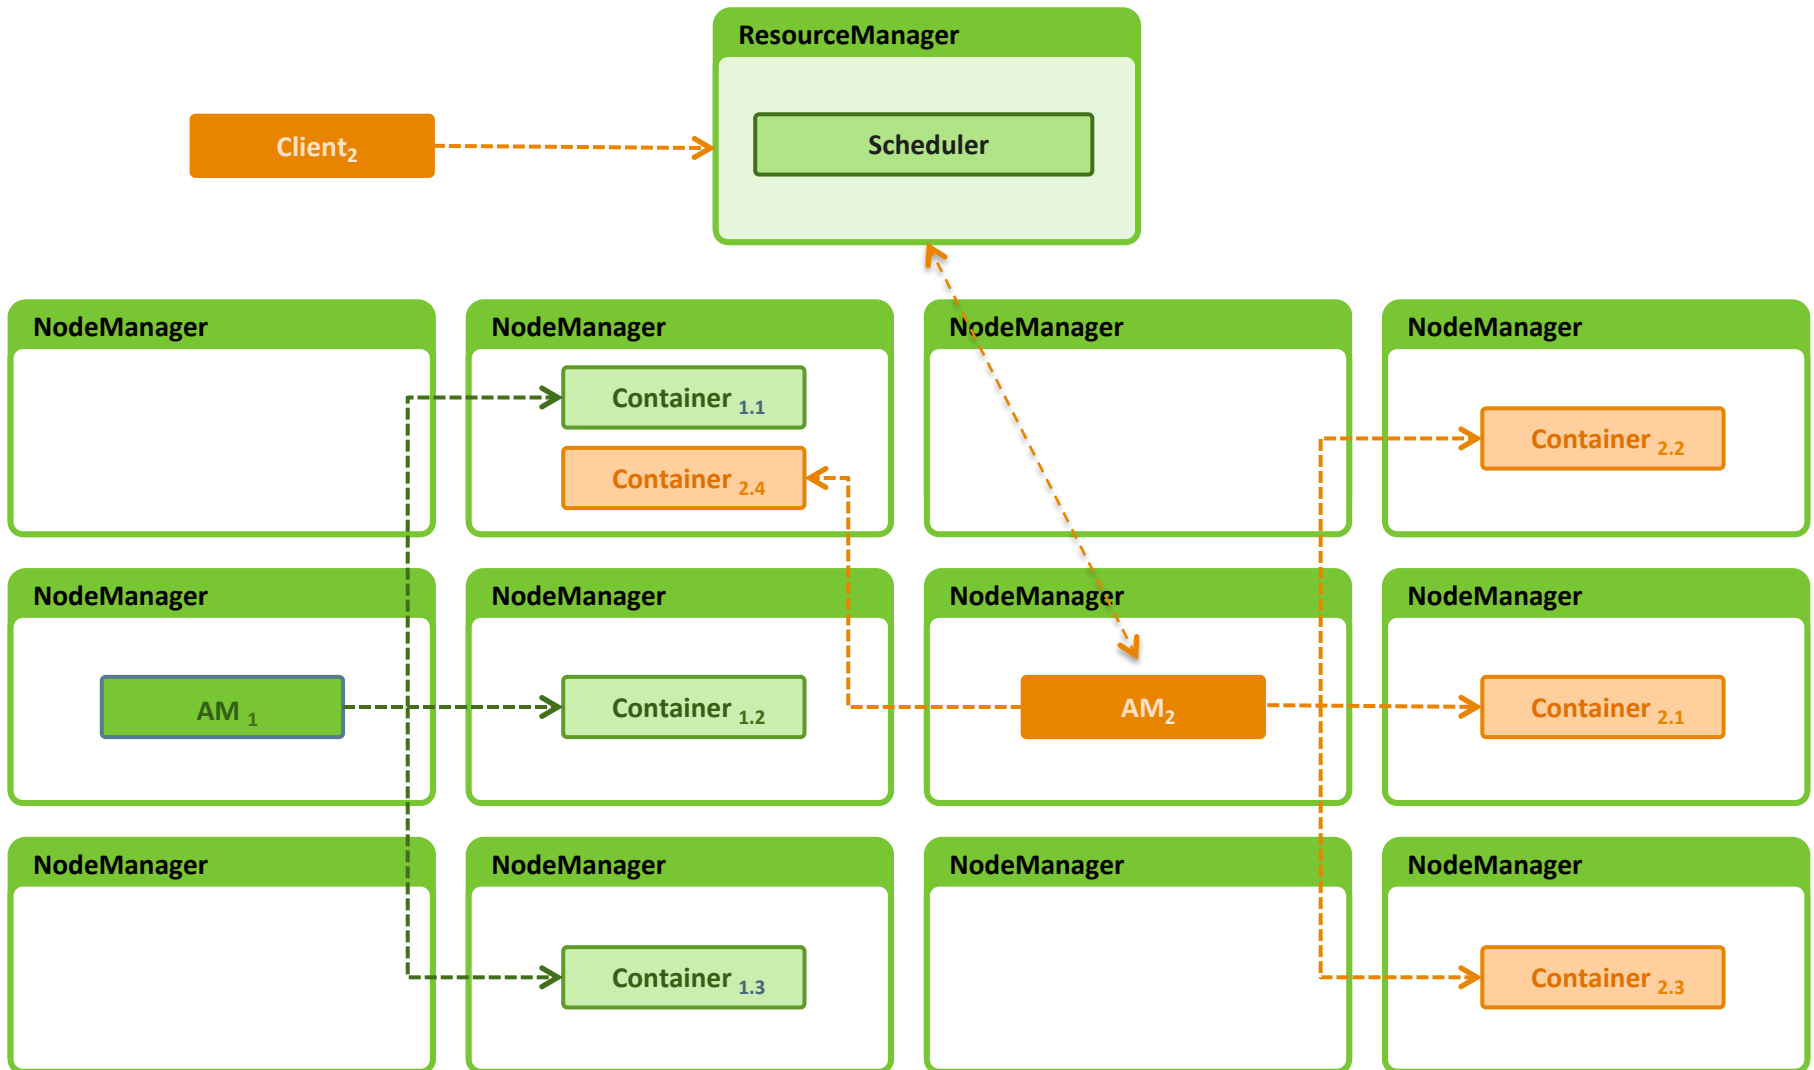
- **Application Master**

- Per-application
- Manages application scheduling and task execution
- E.g. MapReduce Application Master

Design Centre

- **Split up the two major functions of JobTracker**
 - Cluster resource management
 - Application life-cycle management
- **MapReduce becomes user-land library**

YARN Architecture - Walkthrough



Review - Benefits of YARN

- 1. Scale**
- 2. New Programming Models & Services**
- 3. Improved cluster utilization**
- 4. Agility**
- 5. Beyond Java**

Agenda

- Why YARN
- YARN Architecture and Concepts
- **Building applications on YARN**
- Next Steps

YARN Applications

- **Data processing applications and services**

- Online Serving – HOYA (HBase on YARN)
- Real-time event processing – Storm, S4, other commercial platforms
- Tez – Generic framework to run a complex DAG
- MPI: OpenMPI, MPICH2
- Master-Worker
- Machine Learning: Spark
- Graph processing: Giraph
- Enabled by allowing the use of paradigm-specific application master

Run all on the same Hadoop cluster!

YARN – Implementing Applications

- **What APIs do I need to use?**
 - Only three *protocols*
 - Client to ResourceManager
 - Application submission
 - ApplicationMaster to ResourceManager
 - Container allocation
 - ApplicationMaster to NodeManager
 - Container launch
 - Use client libraries for all 3 actions
 - Module `yarn-client`
 - Provides both synchronous and asynchronous libraries
 - Use 3rd party like Weave
 - <http://continuity.github.io/weave/>

YARN – Implementing Applications

- **What do I need to do?**
 - Write a submission Client
 - Write an ApplicationMaster (well copy-paste)
 - *DistributedShell is the new WordCount*
 - Get containers, run whatever you want!

YARN – Implementing Applications

- **What else do I need to *know*?**
 - Resource Allocation & Usage
 - ResourceRequest
 - Container
 - ContainerLaunchContext
 - LocalResource
 - ApplicationMaster
 - ApplicationId
 - ApplicationAttemptId
 - ApplicationSubmissionContext

YARN – Resource Allocation & Usage

- **ResourceRequest**

- Fine-grained resource *ask* to the ResourceManager
- Ask for a specific amount of resources (memory, cpu etc.) on a specific machine or rack
- Use special value of * for resource name for *any* machine

ResourceRequest
priority
resourceName
capability
numContainers

YARN – Resource Allocation & Usage

- **ResourceRequest**

priority	capability	resourceName	numContainers
0	<2gb, 1 core>	host01	1
		rack0	1
		*	1
1	<4gb, 1 core>	*	1

YARN – Resource Allocation & Usage

- **Container**

- The basic unit of allocation in YARN
- The *result* of the ResourceRequest provided by ResourceManager to the ApplicationMaster
- *A specific amount of resources (cpu, memory etc.) on a specific machine*

Container
containerId
resourceName
capability
tokens

YARN – Resource Allocation & Usage

- **ContainerLaunchContext**

- The context provided by ApplicationMaster to NodeManager to launch the Container
- *Complete specification for a process*
- LocalResource used to specify container binary and dependencies
 - NodeManager responsible for downloading from shared namespace (typically HDFS)

ContainerLaunchContext
container
commands
environment
localResources



LocalResource
uri
type

YARN - ApplicationMaster

- **ApplicationMaster**

- Per-application controller aka *container_0*
- Parent for all containers of the application
 - ApplicationMaster negotiates all its containers from ResourceManager
- ApplicationMaster container is child of ResourceManager
 - Think *init* process in Unix
 - RM restarts the ApplicationMaster *attempt* if required (unique `ApplicationAttemptId`)
- Code for application is submitted along with Application itself

YARN - ApplicationMaster

- **ApplicationMaster**

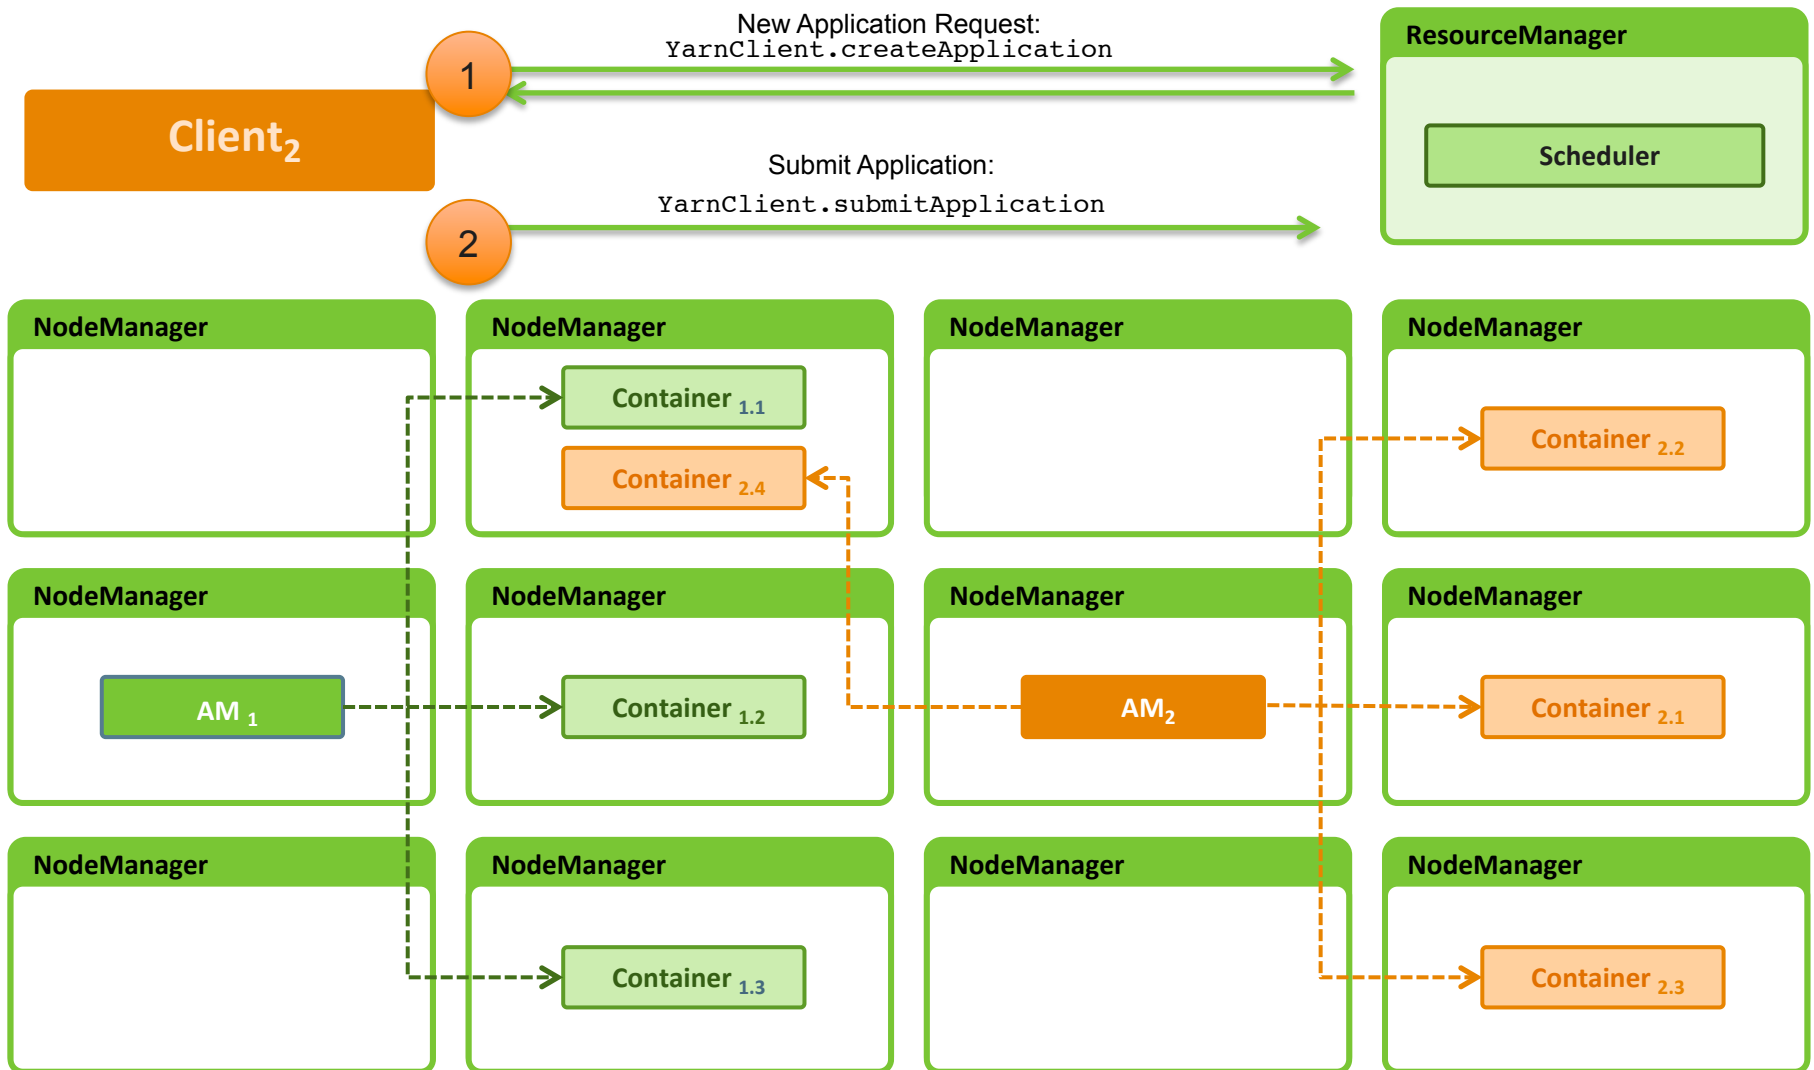
- `ApplicationSubmissionContext` is the complete specification of the `ApplicationMaster`, provided by Client
- `ResourceManager` responsible for *allocating* and *launching* `ApplicationMaster` container

ApplicationSubmissionContext	
	resourceRequest
containerLaunchContext	
	appName
	queue

YARN Application API - Overview

- **hadoop-yarn-client** module
- **YarnClient** is submission client api
- **Both synchronous & asynchronous APIs for resource allocation and container start/stop**
- **Synchronous API**
 - `AMRMClient`
 - `AMNMClient`
- **Asynchronous API**
 - `AMRMClientAsync`
 - `AMNMClientAsync`

YARN Application API – The Client

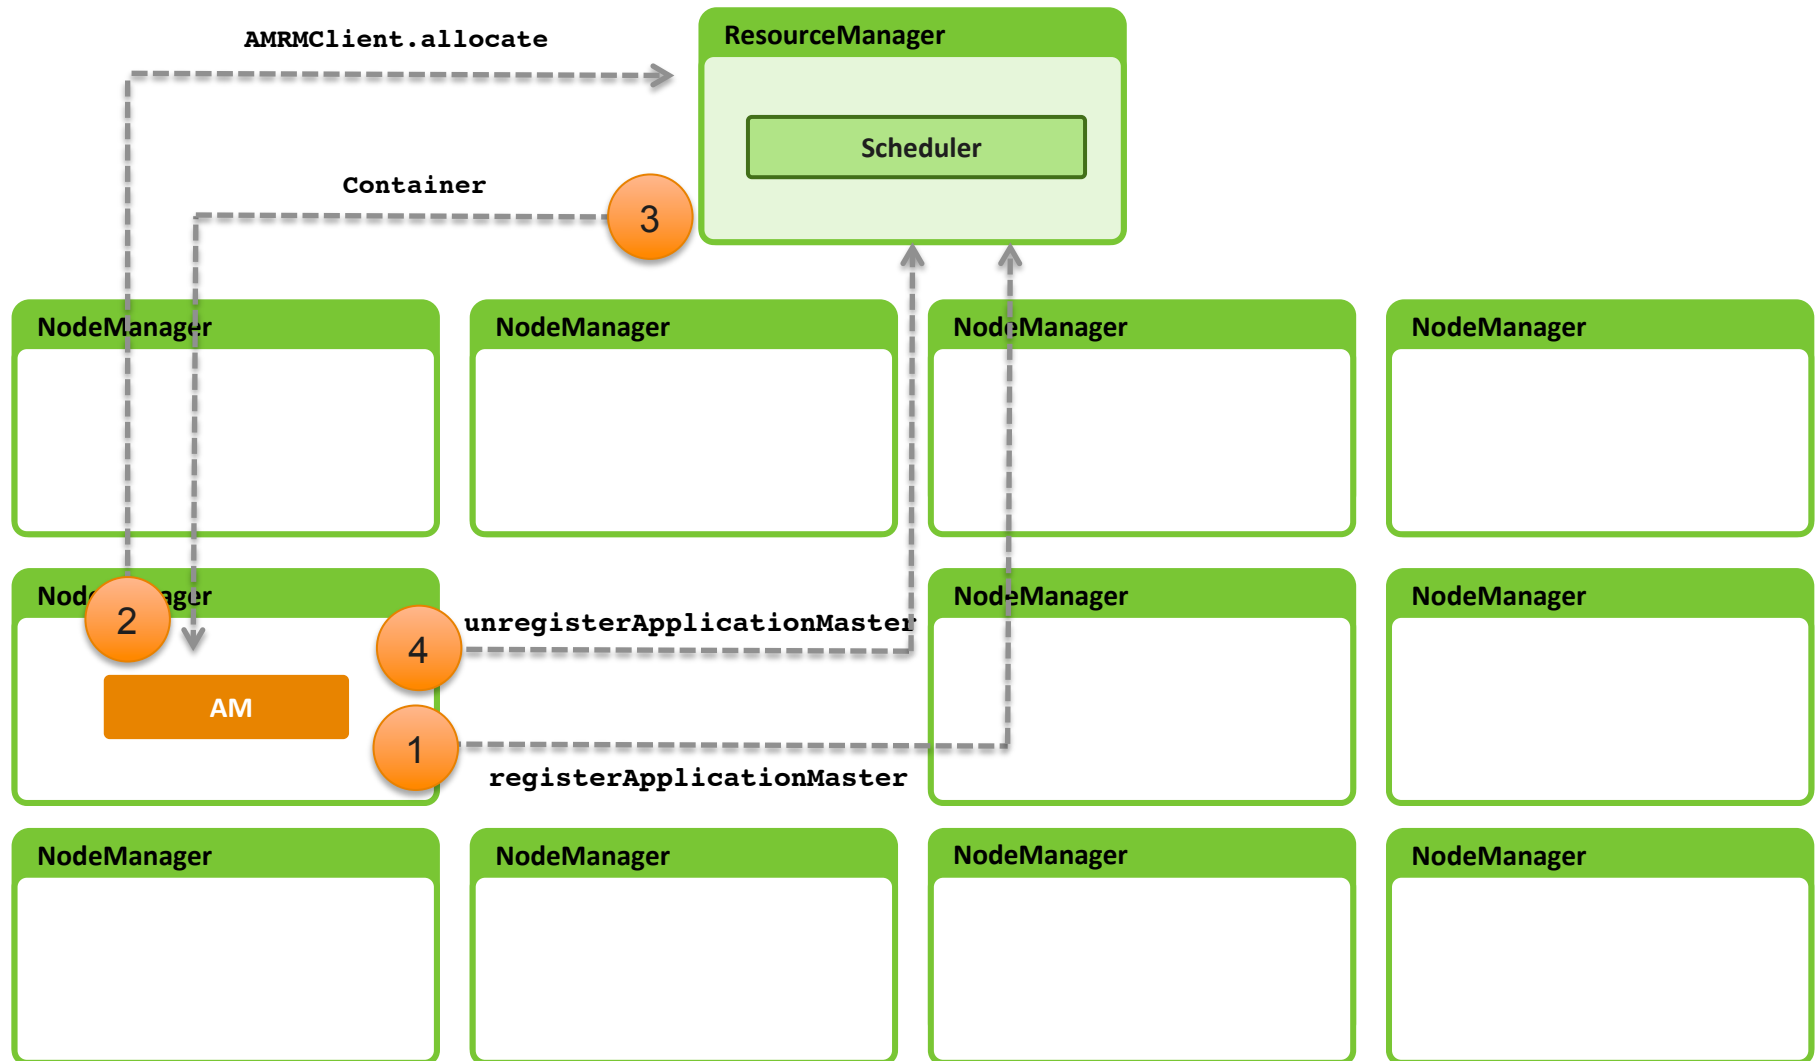


YARN Application API – The Client

- **YarnClient**

- `createApplication` to create application
- `submitApplication` to start application
 - Application developer needs to provide `ApplicationSubmissionContext`
- APIs to get other information from `ResourceManager`
 - `getAllQueues`
 - `getApplications`
 - `getNodeReports`
- APIs to manipulate submitted application e.g. `killApplication`

YARN Application API – Resource Allocation



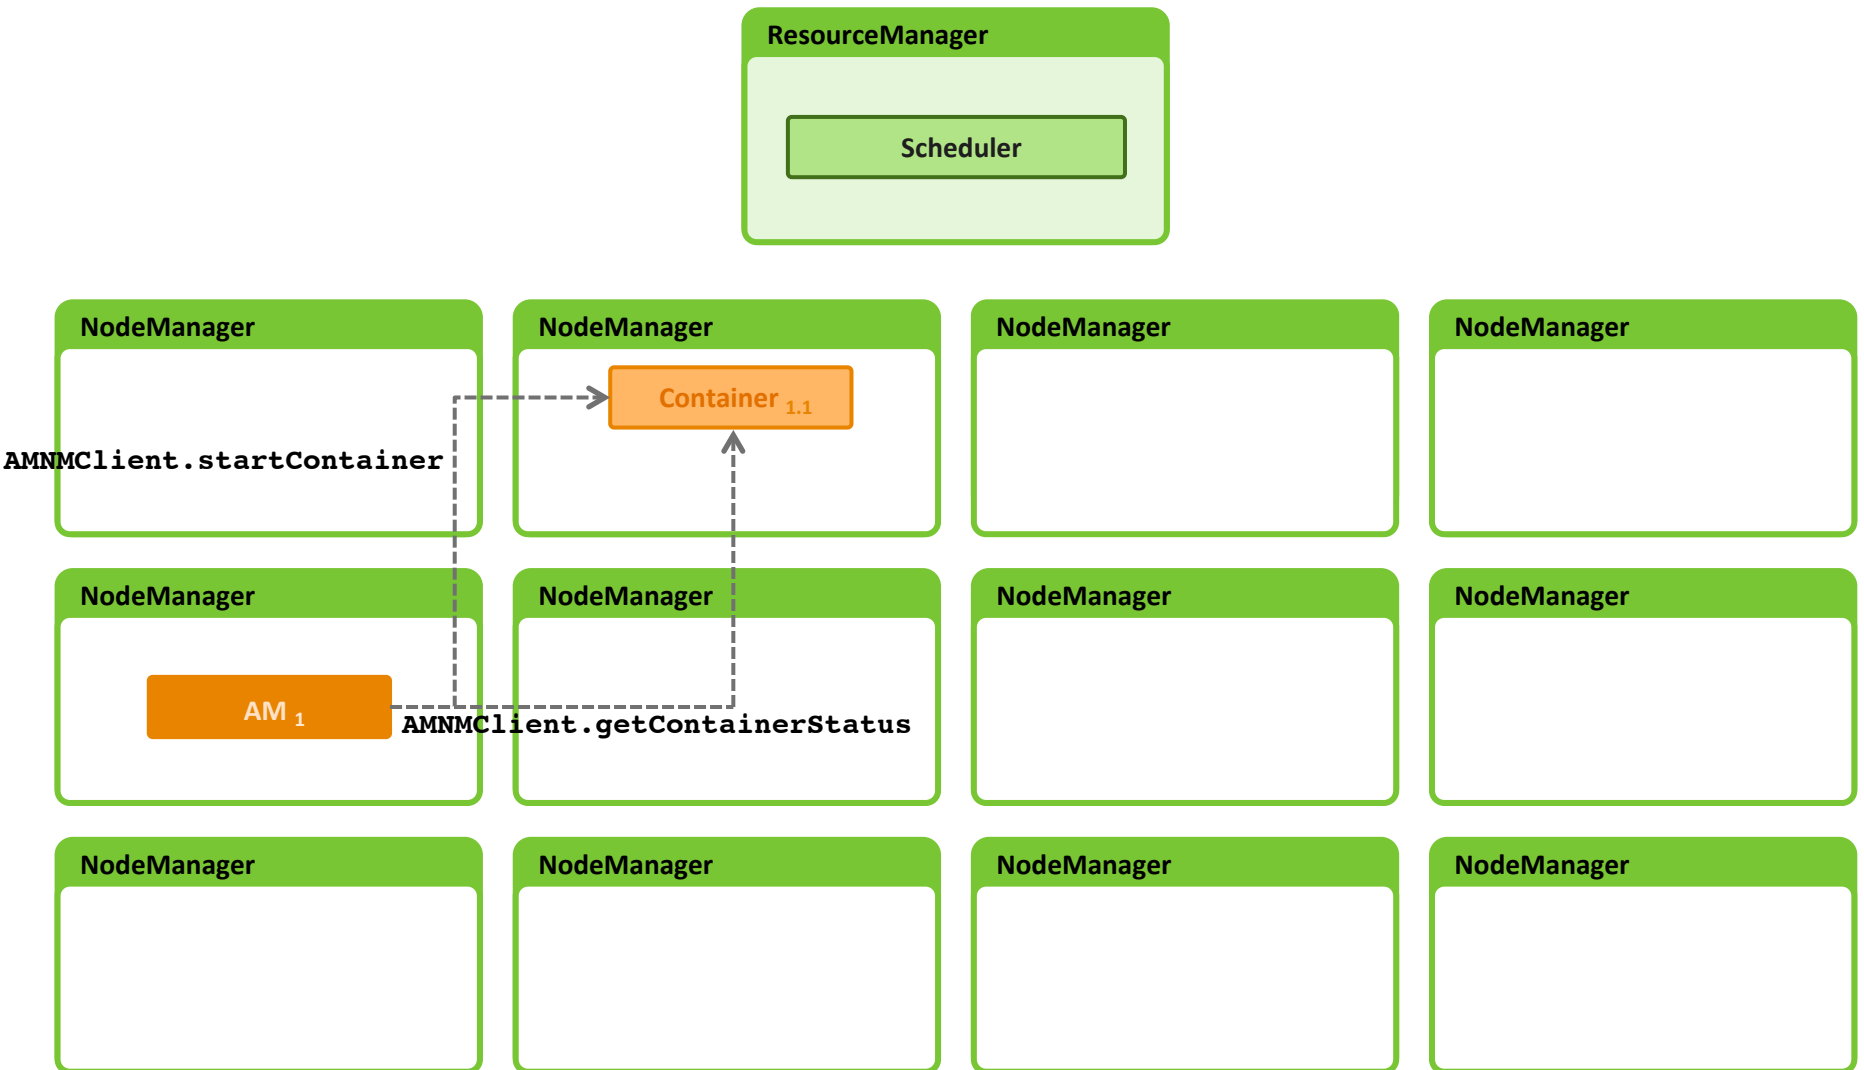
YARN Application API – Resource Allocation

- **AMRMClient - Synchronous API for ApplicationMaster to interact with ResourceManager**
 - Prologue / epilogue – `registerApplicationMaster` / `unregisterApplicationMaster`
 - Resource negotiation with ResourceManager
 - Internal book-keeping - `addContainerRequest` / `removeContainerRequest` / `releaseAssignedContainer`
 - Main API – `allocate`
 - Helper APIs for cluster information
 - `getAvailableResources`
 - `getClusterNodeCount`

YARN Application API – Resource Allocation

- **AMRMClientAsync - Asynchronous API for ApplicationMaster**
 - Extension of `AMRMClient` to provide asynchronous `CallbackHandler`
 - Callbacks make it easier to build mental model of interaction with `ResourceManager` for the application developer
 - `onContainersAllocated`
 - `onContainersCompleted`
 - `onNodesUpdated`
 - `onError`
 - `onShutdownRequest`

YARN Application API – Using Resources



YARN Application API – Using Resources

- **AMNMClient** - Synchronous API for ApplicationMaster to launch / stop containers at NodeManager
 - Simple (trivial) APIs
 - startContainer
 - stopContainer
 - getContainerStatus

YARN Application API – Using Resources

- **AMNMClient - Asynchronous API for ApplicationMaster to launch / stop containers at NodeManager**
 - Simple (trivial) APIs
 - startContainerAsync
 - stopContainerAsync
 - getContainerStatusAsync
 - CallbackHandler to make it easier to build mental model of interaction with NodeManager for the application developer
 - onContainerStarted
 - onContainerStopped
 - onStartContainerError
 - onContainerStatusReceived

YARN Application API - Development

- **Un-Managed Mode for ApplicationMaster**

- Run ApplicationMaster on development machine rather than in-cluster
 - No submission client
- `hadoop-yarn-applications-unmanaged-am-launcher`
- Easier to step through debugger, browse logs etc.

```
$ bin/hadoop jar hadoop-yarn-applications-unmanaged-am-launcher.jar \  
  Client \  
  -jar my-application-master.jar \  
  -cmd 'java MyApplicationMaster <args>'
```

Sample YARN Application – DistributedShell

- **Overview**

- YARN application to run n copies for a Shell command
- Simplest example of a YARN application – get n containers and run a specific Unix command

```
$ bin/hadoop jar hadoop-yarn-applications-distributedshell.jar \  
    org.apache.hadoop.yarn.applications.distributedshell.Client \  
    -shell_command '/bin/date' \  
    -num_containers <n>
```

Code: <https://github.com/hortonworks/simple-yarn-app>

Sample YARN Application – DistributedShell

- **Code Overview**

- User submits application to ResourceManager via `org.apache.hadoop.yarn.applications.distributedshell.Client`
 - Client provides `ApplicationSubmissionContext` to the ResourceManager
- It is responsibility of `org.apache.hadoop.yarn.applications.distributedshell.ApplicationMaster` to negotiate n containers
 - ApplicationMaster launches containers with the user-specified *command* as `ContainerLaunchContext.commands`

Sample YARN Application – DistributedShell

- **Client – Code Walkthrough**

- hadoop-yarn-client module
- Steps:
 - YarnClient.createApplication
 - Specify ApplicationSubmissionContext, in particular, ContainerLaunchContext with commands, and other key pieces such as resource capability for ApplicationMaster container and queue, appName, appType etc.
 - YarnClient.submitApplication

```
1.    // Create yarnClient
2.    YarnClient yarnClient = YarnClient.createYarnClient();
3.    yarnClient.init(new Configuration());
4.    yarnClient.start();
5.
6.    // Create application via yarnClient
7.    YarnClientApplication app = yarnClient.createApplication();
8.
```

Sample YARN Application – DistributedShell

• Client – Code Walkthrough

```
9.      // Set up the container launch context for the application master
10.     ContainerLaunchContext amContainer =
11.         Records.newRecord(ContainerLaunchContext.class);
12.     List<String> command = new List<String>();
13.     commands.add("$JAVA_HOME/bin/java");
14.     commands.add("-Xmx256M");
15.     commands.add(
16.         "org.apache.hadoop.yarn.applications.distributedshell.ApplicationMaster");
17.     commands.add("--container_memory 1024");
18.     commands.add("--container_cores 1");
19.     commands.add("--num_containers 3");
20.     amContainer.setCommands(commands);
21.
22.     // Set up resource type requirements for ApplicationMaster
23.     Resource capability = Records.newRecord(Resource.class);
24.     capability.setMemory(256);
25.     capability.setVirtualCores(2);
26.
```

Command to launch ApplicationMaster process

Resources required for ApplicationMaster container

Sample YARN Application – DistributedShell

- **Client – Code Walkthrough**

```
27. // Finally, set-up ApplicationSubmissionContext for the application
28. ApplicationSubmissionContext appContext =
29.     app.getApplicationSubmissionContext();
30. appContext.setQueue("my-queue"); // queue
31. appContext.setAMContainerSpec(amContainer);
32. appContext.setResource(capability);
33. appContext.setApplicationName("my-app"); // application name
34. appContext.setApplicationType("DISTRIBUTED_SHELL"); // application type
35.
36. // Submit application
37. yarnClient.submitApplication(appContext);
```

ApplicationSubmissionContext
for
ApplicationMaster

Submit application to
ResourceManager

Sample YARN Application – DistributedShell

- **ApplicationMaster – Code Walkthrough**

- Again, `hadoop-yarn-client` module
- Steps:
 - `AMRMClient.registerApplication`
 - Negotiate containers from ResourceManager by providing `ContainerRequest` to `AMRMClient.addContainerRequest`
 - Take the resultant `Container` returned via subsequent call to `AMRMClient.allocate`, build `ContainerLaunchContext` with `Container` and commands, then launch them using `AMNMClient.launchContainer`
 - Use `LocalResources` to specify software/configuration dependencies for each worker container
 - Wait till done...
`AllocateResponse.getCompletedContainersStatuses` from subsequent calls to `AMRMClient.allocate`
 - `AMRMClient.unregisterApplication`

Sample YARN Application – DistributedShell

• ApplicationMaster – Code Walkthrough

```
1.      // Initialize clients to ResourceManager and NodeManagers
2.      Configuration conf = new Configuration();
3.
4.      AMRMClient rmClient = AMRMClientAsync.createAMRMClient();
5.      rmClient.init(conf);
6.      rmClient.start();
7.
8.      NMClient nmClient = NMClient.createNMClient();
9.      nmClientAsync.init(conf);
10.     nmClientAsync.start();
11.
12.     // Register with ResourceManager
13.     rmClient.registerApplicationMaster("", 0, "");
```

Initialize clients to
ResourceManager and
NodeManagers

Register with
ResourceManager

Sample YARN Application – DistributedShell

• ApplicationMaster – Code Walkthrough

```
15.    // Priority for worker containers - priorities are intra-application
16.    Priority priority = Records.newRecord(Priority.class);
17.    priority.setPriority(0);
18.
19.    // Resource requirements for worker containers
20.    Resource capability = Records.newRecord(Resource.class);
21.    capability.setMemory(128);
22.    capability.setVirtualCores(1);
23.
24.    // Make container requests to ResourceManager
25.    for (int i = 0; i < n; ++i) {
26.        ContainerRequest containerAsk = new ContainerRequest(capability, null, null,
priority);
27.        rmClient.addContainerRequest(containerAsk);
28.    }
```

Setup requirements for
worker containers

Make resource requests
to ResourceManager

Sample YARN Application – DistributedShell

• ApplicationMaster – Code Walkthrough

```
30. // Obtain allocated containers and launch
31. int allocatedContainers = 0;
32. while (allocatedContainers < n) {
33.     AllocateResponse response = rmClient.allocate(0);
34.     for (Container container : response.getAllocatedContainers()) {
35.         ++allocatedContainers;
36.
37.         // Launch container by create ContainerLaunchContext
38.         ContainerLaunchContext ctx = Records.newRecord(ContainerLaunchContext.class);
39.         ctx.setCommands(Collections.singletonList("/bin/date"));
40.         nmClient.startContainer(container, ctx);
41.     }
42.     Thread.sleep(100);
43. }
```

Setup requirements for
worker containers

Make resource requests
to ResourceManager

Sample YARN Application – DistributedShell

• ApplicationMaster – Code Walkthrough

```
4 // Now wait for containers to complete
4 int completedContainers = 0;
4 while (completedContainers < n) {
4     AllocateResponse response = rmClient.allocate(completedContainers/n);
4     for (ContainerStatus status : response.getCompletedContainersStatuses()) {
5         if (status.getExitStatus() == 0) {
5             ++completedContainers;
5         }
5     }
4     Thread.sleep(100);
5 }
5
5 // Un-register with ResourceManager
5 rmClient.unregisterApplicationMaster(SUCCEEDED, "", "");
```

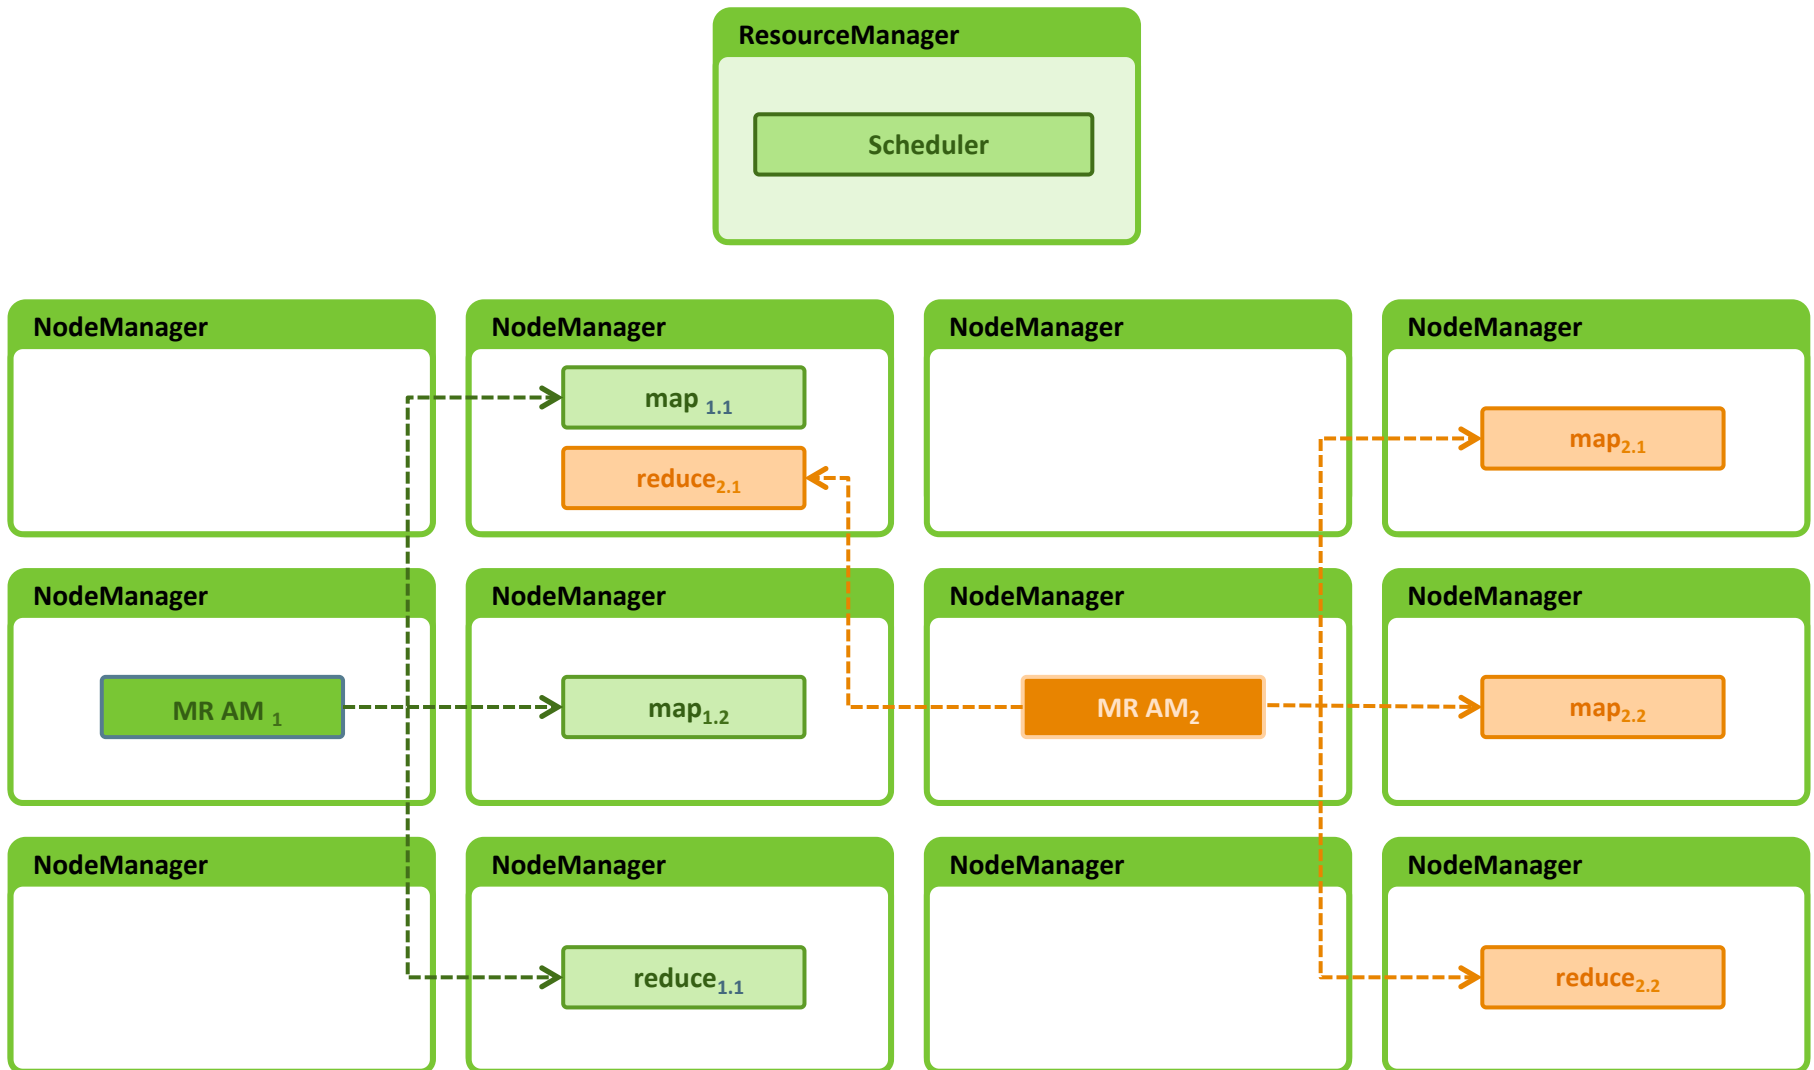
Wait for containers to
complete successfully

Un-register with
ResourceManager

Apache Hadoop MapReduce on YARN

- **Original use-case**
- **Most complex application to build**
 - Data-locality
 - Fault tolerance
 - ApplicationMaster recovery: Check point to HDFS
 - Intra-application Priorities: Maps v/s Reduces
 - Needed complex feedback mechanism from ResourceManager
 - Security
 - Isolation
- **Binary compatible with Apache Hadoop 1.x**

Apache Hadoop MapReduce on YARN




Apache Tez on YARN


- **Replaces MapReduce as primitive for Pig, Hive, Cascading etc.**

- Smaller latency for interactive queries
- Higher throughput for batch queries

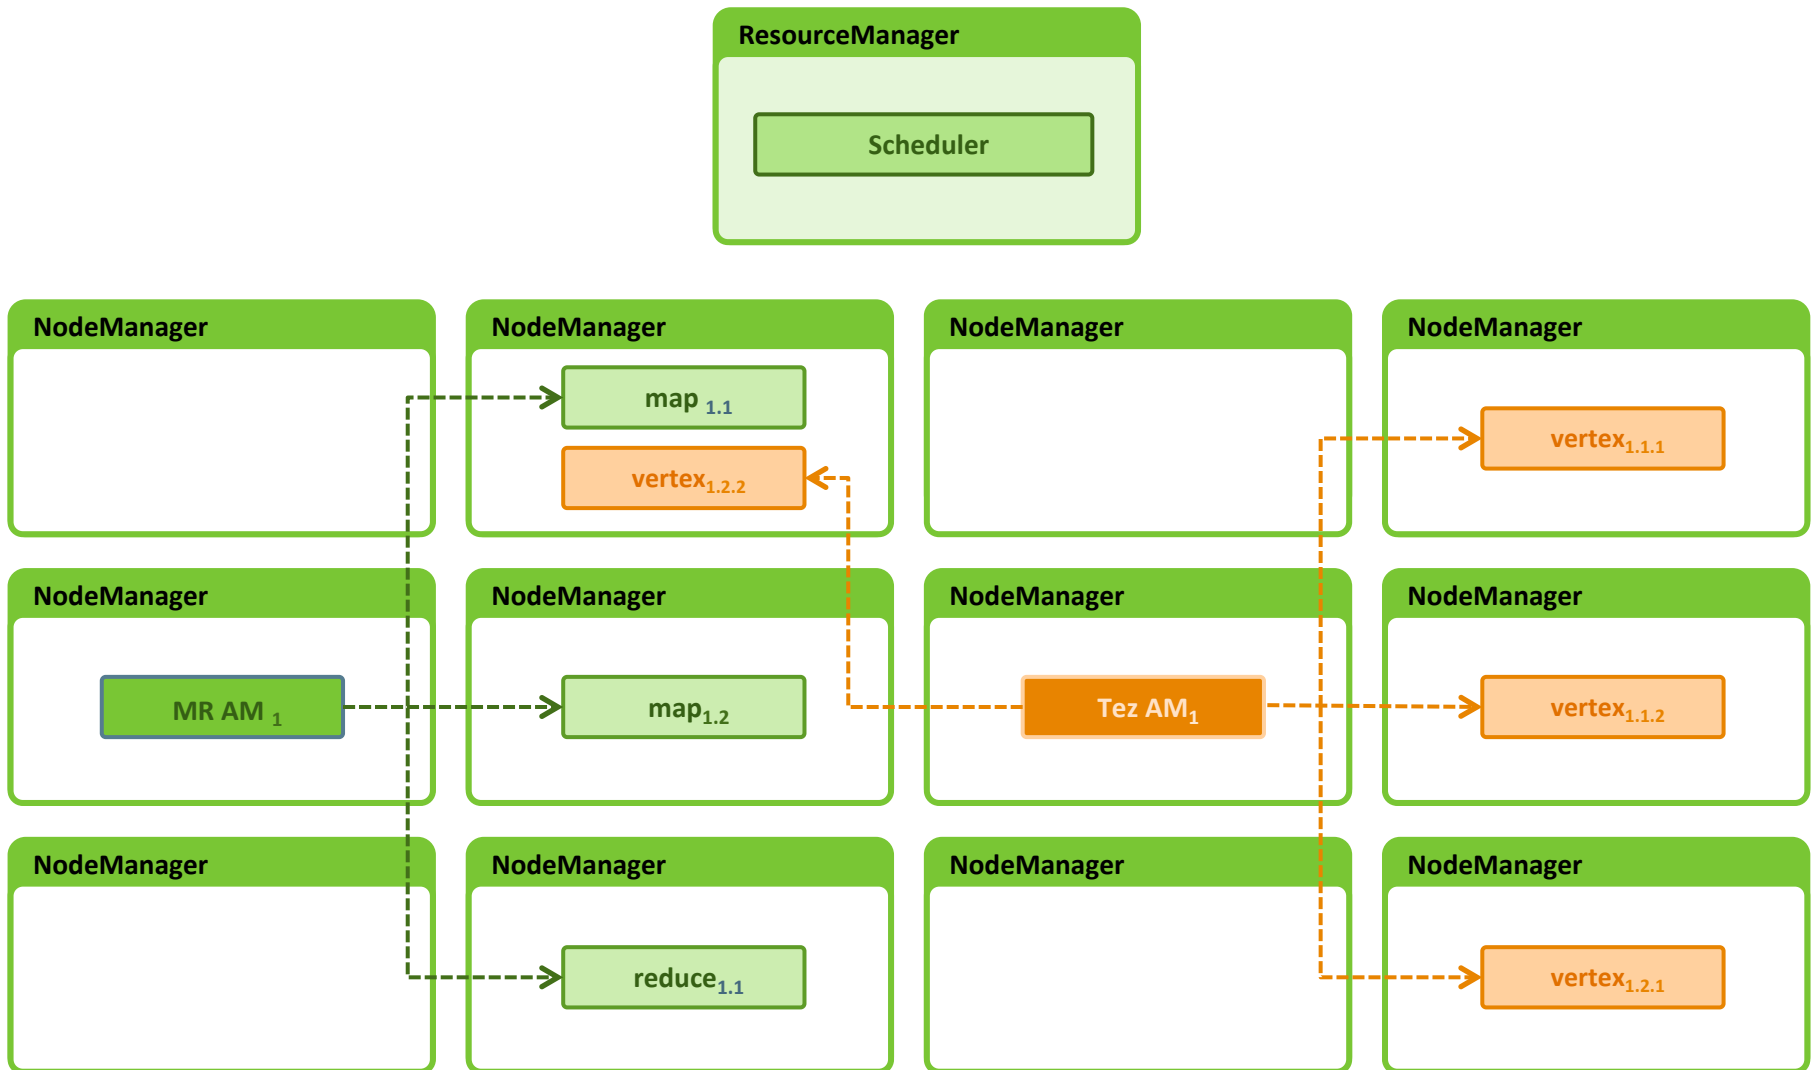
– <http://incubator.apache.org/projects/tez.html>



```
SELECT a.x, AVERAGE(b.y) AS avg
FROM a JOIN b ON (a.id = b.id)
GROUP BY a
UNION SELECT x, AVERAGE(y) AS AVG
FROM c
GROUP BY x
ORDER BY AVG;
```



Apache Tez on YARN

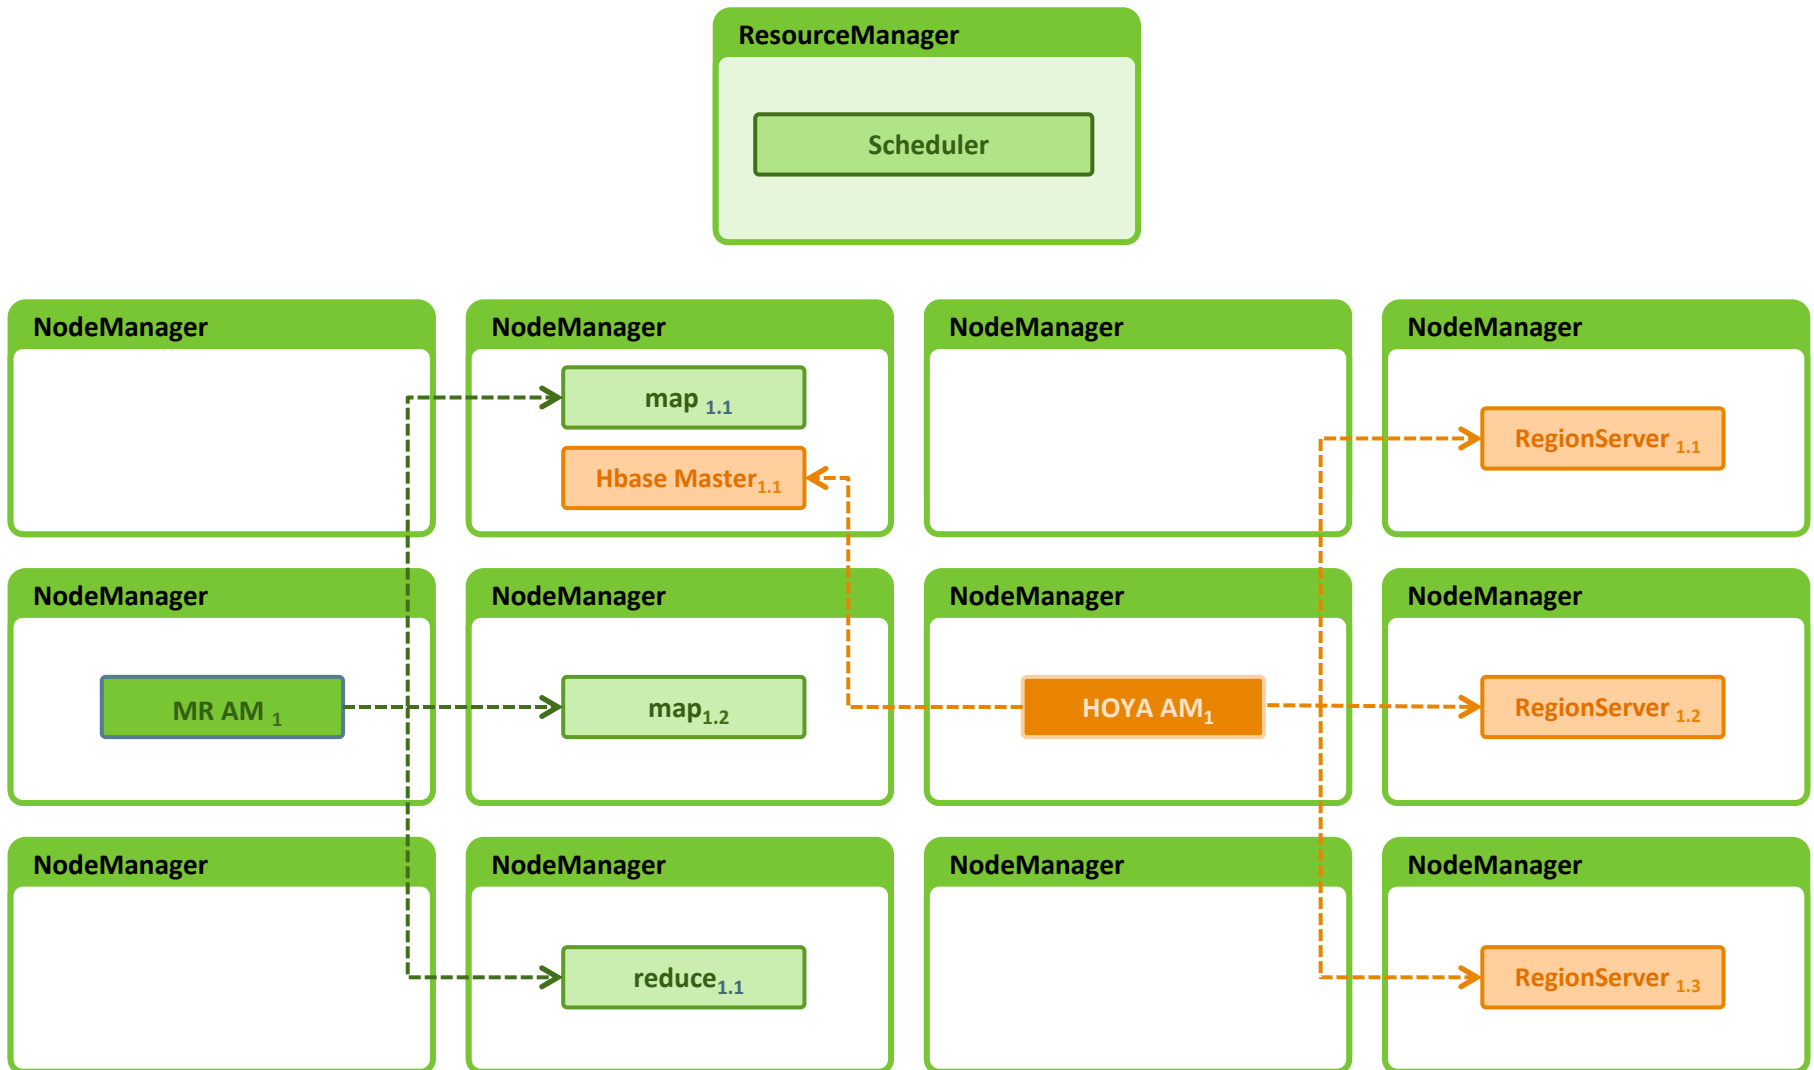


HOYA - Apache HBase on YARN

- **Hoya – Apache HBase becomes *user-level* application**
- **Use cases**
 - Small HBase cluster in large YARN cluster
 - Dynamic HBase clusters
 - Transient/intermittent clusters for workflows
- **APIs to create, start, stop & delete HBase clusters**
- **Flex cluster size: increase/decrease size with load**
- **Recover from Region Server loss with new container.**

Code: <https://github.com/hortonworks/hoya>

HOYA - Apache HBase on YARN



HOYA - Highlights

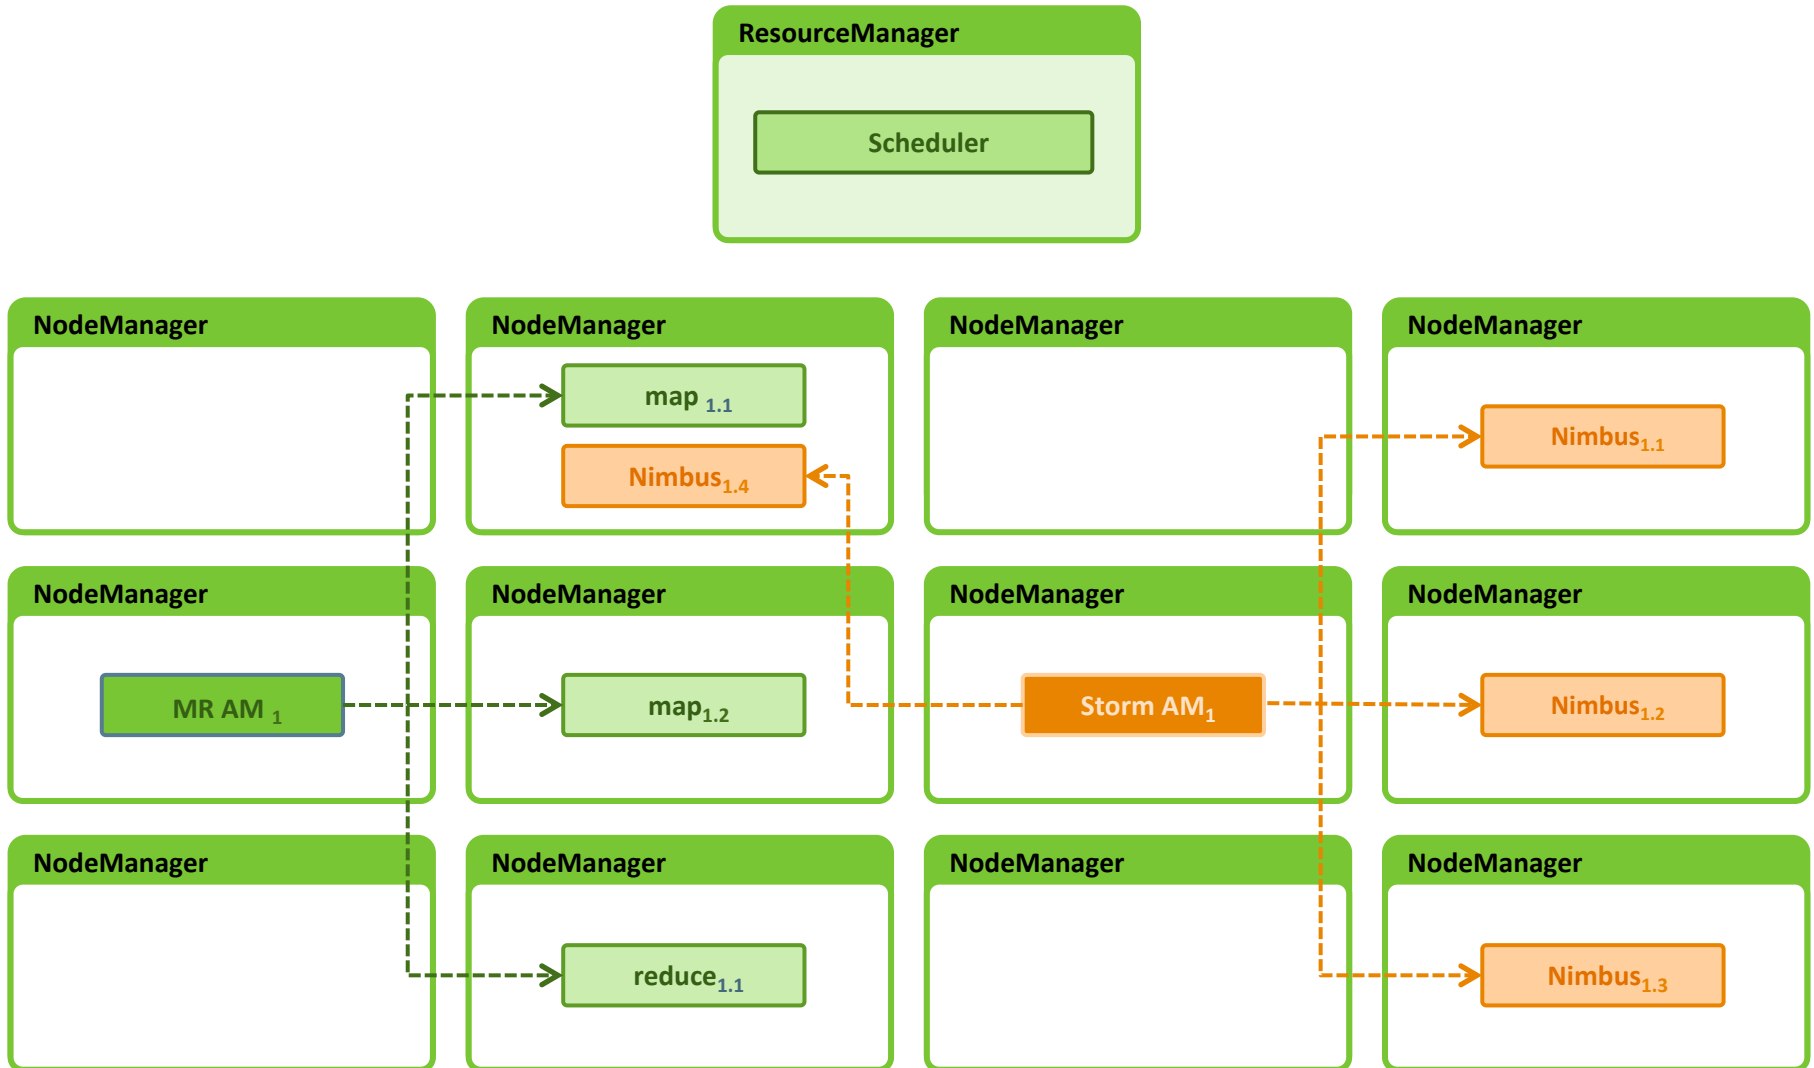
- **Cluster specification stored as JSON in HDFS**
- **Config directory cached - dynamically patched before pushing up as local resources for Master & RegionServers**
- **HBase tar file stored in HDFS -clusters can use the same/different HBase versions**
- **Handling of cluster flexing is the same code as unplanned container loss.**
- **No Hoya code on RegionServers: client and AM only**

Storm on YARN

- **Ability to deploy multiple Storm clusters on YARN for real-time event processing**
- **Yahoo – Primary contributor**
 - 200+ nodes in production
- **Ability to recover from faulty nodes**
 - Get new containers
- **Auto-scale for load balancing**
 - Get new containers as load increases
 - Release containers as load decreases

Code: <https://github.com/yahoo/storm-yarn>

Storm on YARN



General Architectural Considerations

- **Fault Tolerance**
 - Checkpoint
- **Security**
- **Always-On services**
- **Scheduler features**
 - Whitelist resources
 - Blacklist resources
 - Labels for machines
 - License management

Agenda

- Why YARN
- YARN Architecture and Concepts
- Building applications on YARN
- **Next Steps**

HDP 2.0 Community Preview & YARN Certification Program

Goal: Accelerate # of certified YARN-based solutions

- **HDP 2.0 Community Preview**

- Contains latest community Beta of Apache Hadoop 2.0 & YARN
- Delivered as easy to use Sandbox VM, as well as RPMs and Tarballs
- Enables YARN Cert Program
 - Community & commercial ecosystem to test and certify new and existing YARN-based apps

- **YARN Certification Program**

- More than 14 partners in program at launch
 - *Splunk**
 - *Elastic Search**
 - *Altiscale**
 - *Concurrent**
 - *Microsoft*
 - *Platfora*
 - *Tableau*
 - *(IBM) DataStage*
 - *Informatica*
 - *Karmasphere*
 - *and others*

* Already certified

Forum & Office Hours

- **YARN Forum**

- Community of Hadoop YARN developers
- Focused on collaboration and Q&A
- <http://hortonworks.com/community/forums/forum/yarn>

- **Office Hours**

- <http://www.meetup.com/HSquared-Hadoop-Hortonworks-User-Group/>
- Bi-weekly office hours with Hortonworks engineers & architects
- Every other Thursday, starting on Aug 15th from 4:30 – 5:30 PM (PST)
- At Hortonworks Palo Alto HQ
 - West Bayshore Rd. Palo Alto, CA 94303 USA

Technical Resources


- **ASF hadoop-2.1.0-beta**

- Coming soon!
- <http://hadoop.apache.org/common/releases.html>
- Release Documentation:
<http://hadoop.apache.org/common/docs/r2.1.0-beta>

- **Blogs:**

- <http://hortonworks.com/hadoop/yarn>
- <http://hortonworks.com/blog/category/apache-hadoop/yarn/>
- <http://hortonworks.com/blog/introducing-apache-hadoop-yarn/>
- <http://hortonworks.com/blog/apache-hadoop-yarn-background-and-an-overview/>

What Next?



- Download the Book
- Download the Community preview
- Join the Beta Program
- Participate via the YARN forum
- Come see us during the *YARN Office Hours*
- Follow us... [@hortonworks](https://twitter.com/hortonworks)

Thank You!

<http://hortonworks.com/hadoop/yarn>