

# Tupperware: Containerized Deployment at FB

DockerCon 2014

Aravind Narayanan  
aravindn@fb.com

facebook

# Scale makes everything harder

- Running single instance: easy
- Running at scale in production: messy and complicated

Provision machines

Machine decoms

Distribute binaries

Failover

Geo-distribution

Monitoring

Daemonize process



**Time spent on getting app  
to run in prod**

| Category                                 | Percentage |
|--|------------|
| Time spent on getting app to run in prod | 60%        |
| Time spent on Application Logic          | 40%        |

**Time spent on  
Application Logic**

# Tupperware to the rescue


“This is my binary. Run it on X machines!”

- Engineer is **hands-off**
  - Doesn't need to worry about machines in prod
- Handles **failover**, when machines go bad
- **Efficient** use of infrastructure
- **300,000+** processes, spread over **15,000+** services

# Agenda

1. Architecture
2. Sandboxes
3. Ecosystem
4. Lessons learnt

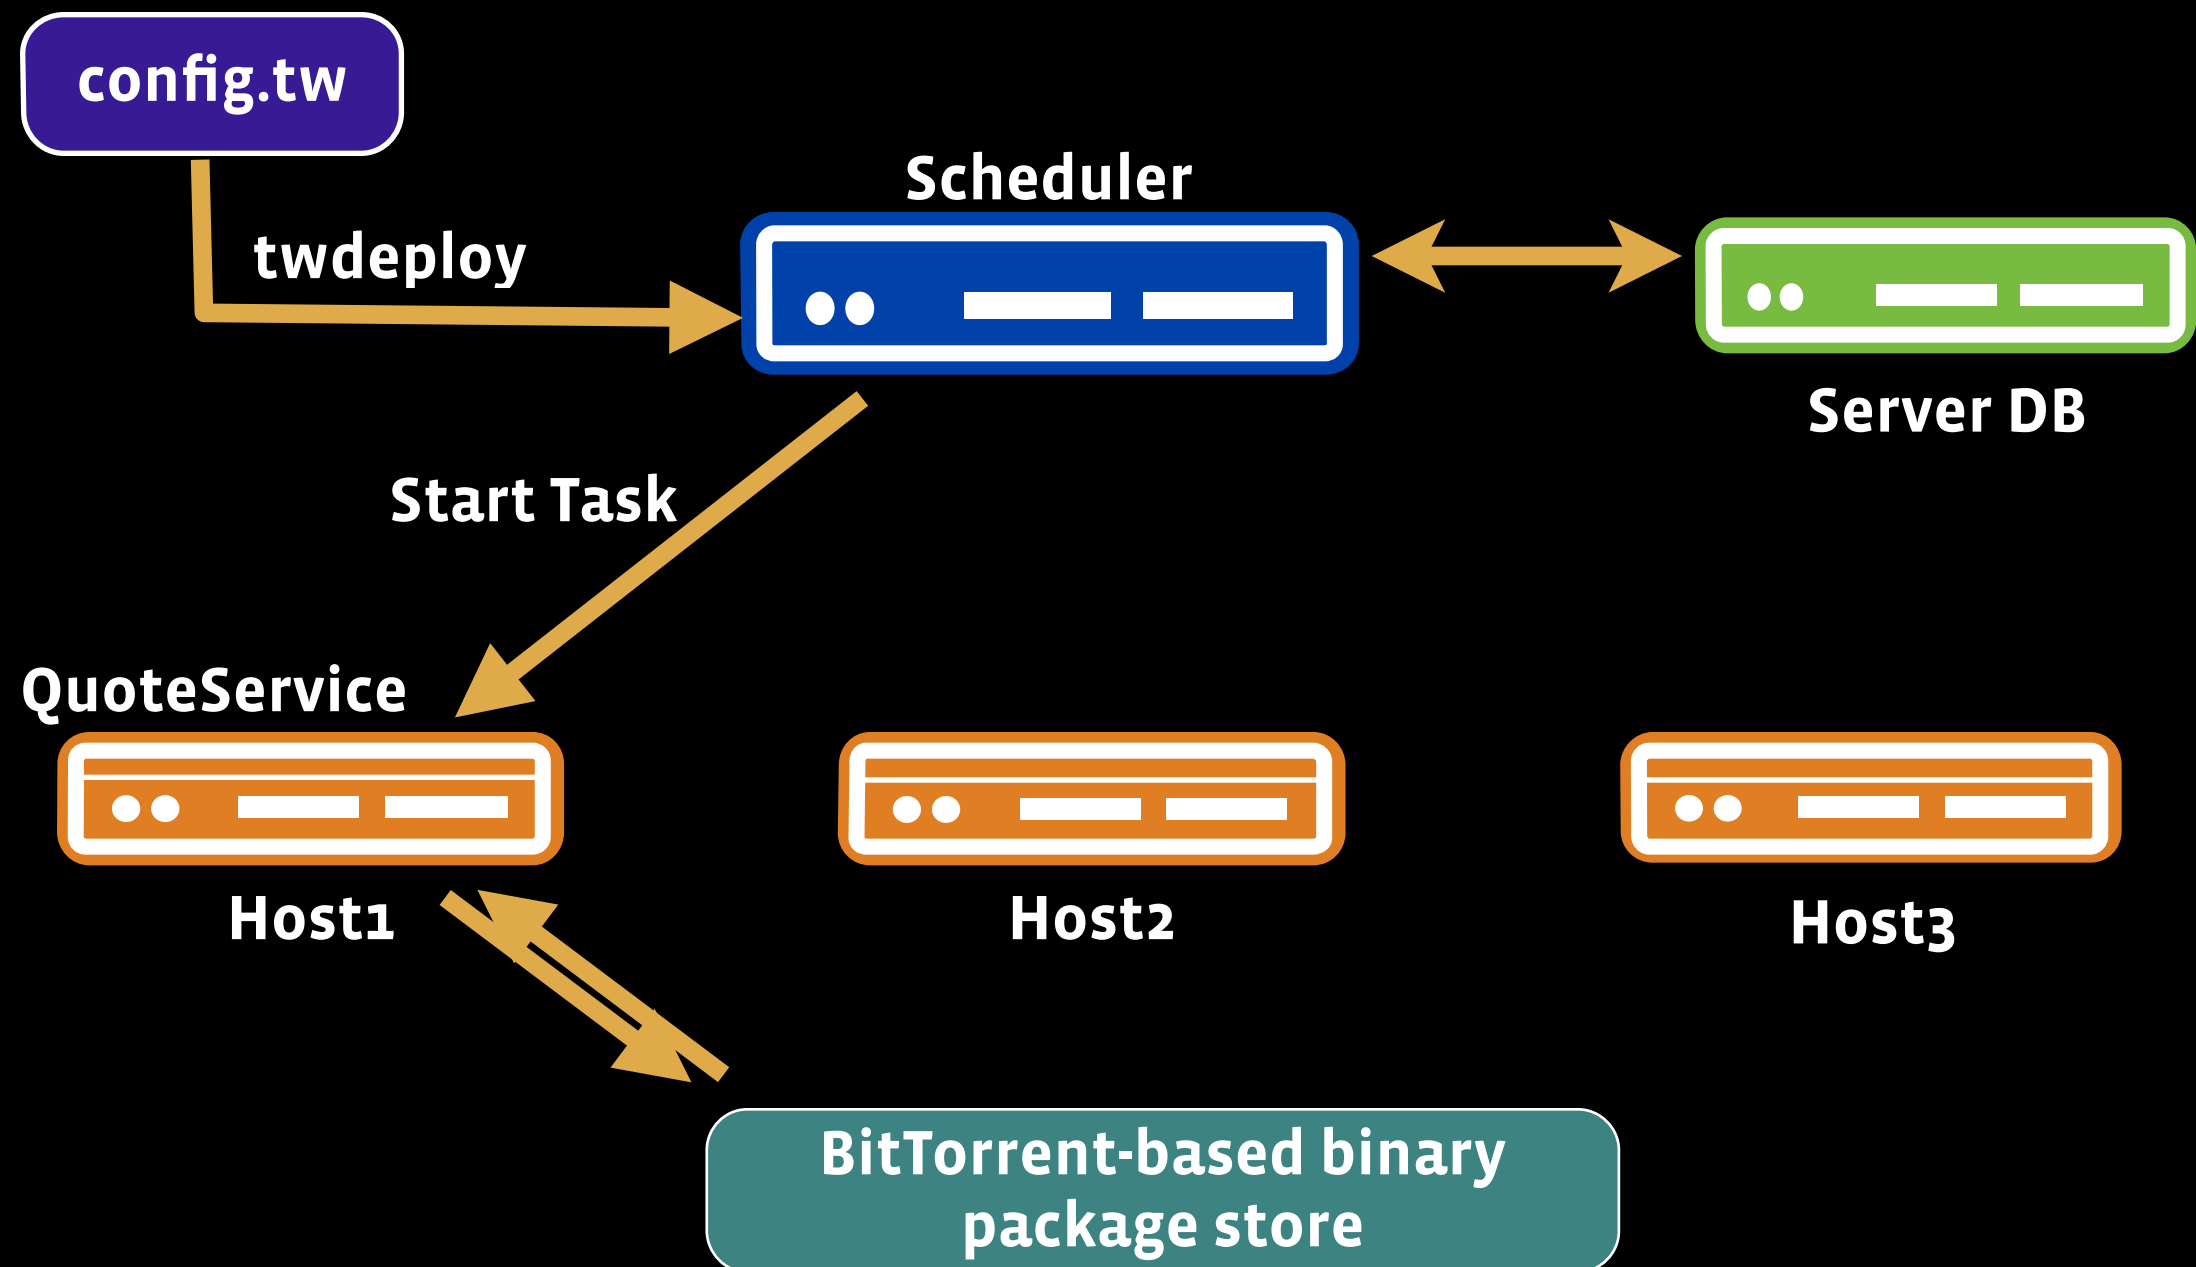
# Facebook Datacenters



# Terminology


- A **DC** has one or more clusters
- A **cluster** has multiple racks
- A **rack** has multiple machines
- A **TW job** is equivalent to a service
- A job has multiple **tasks**, each an instance of the service

# Architecture



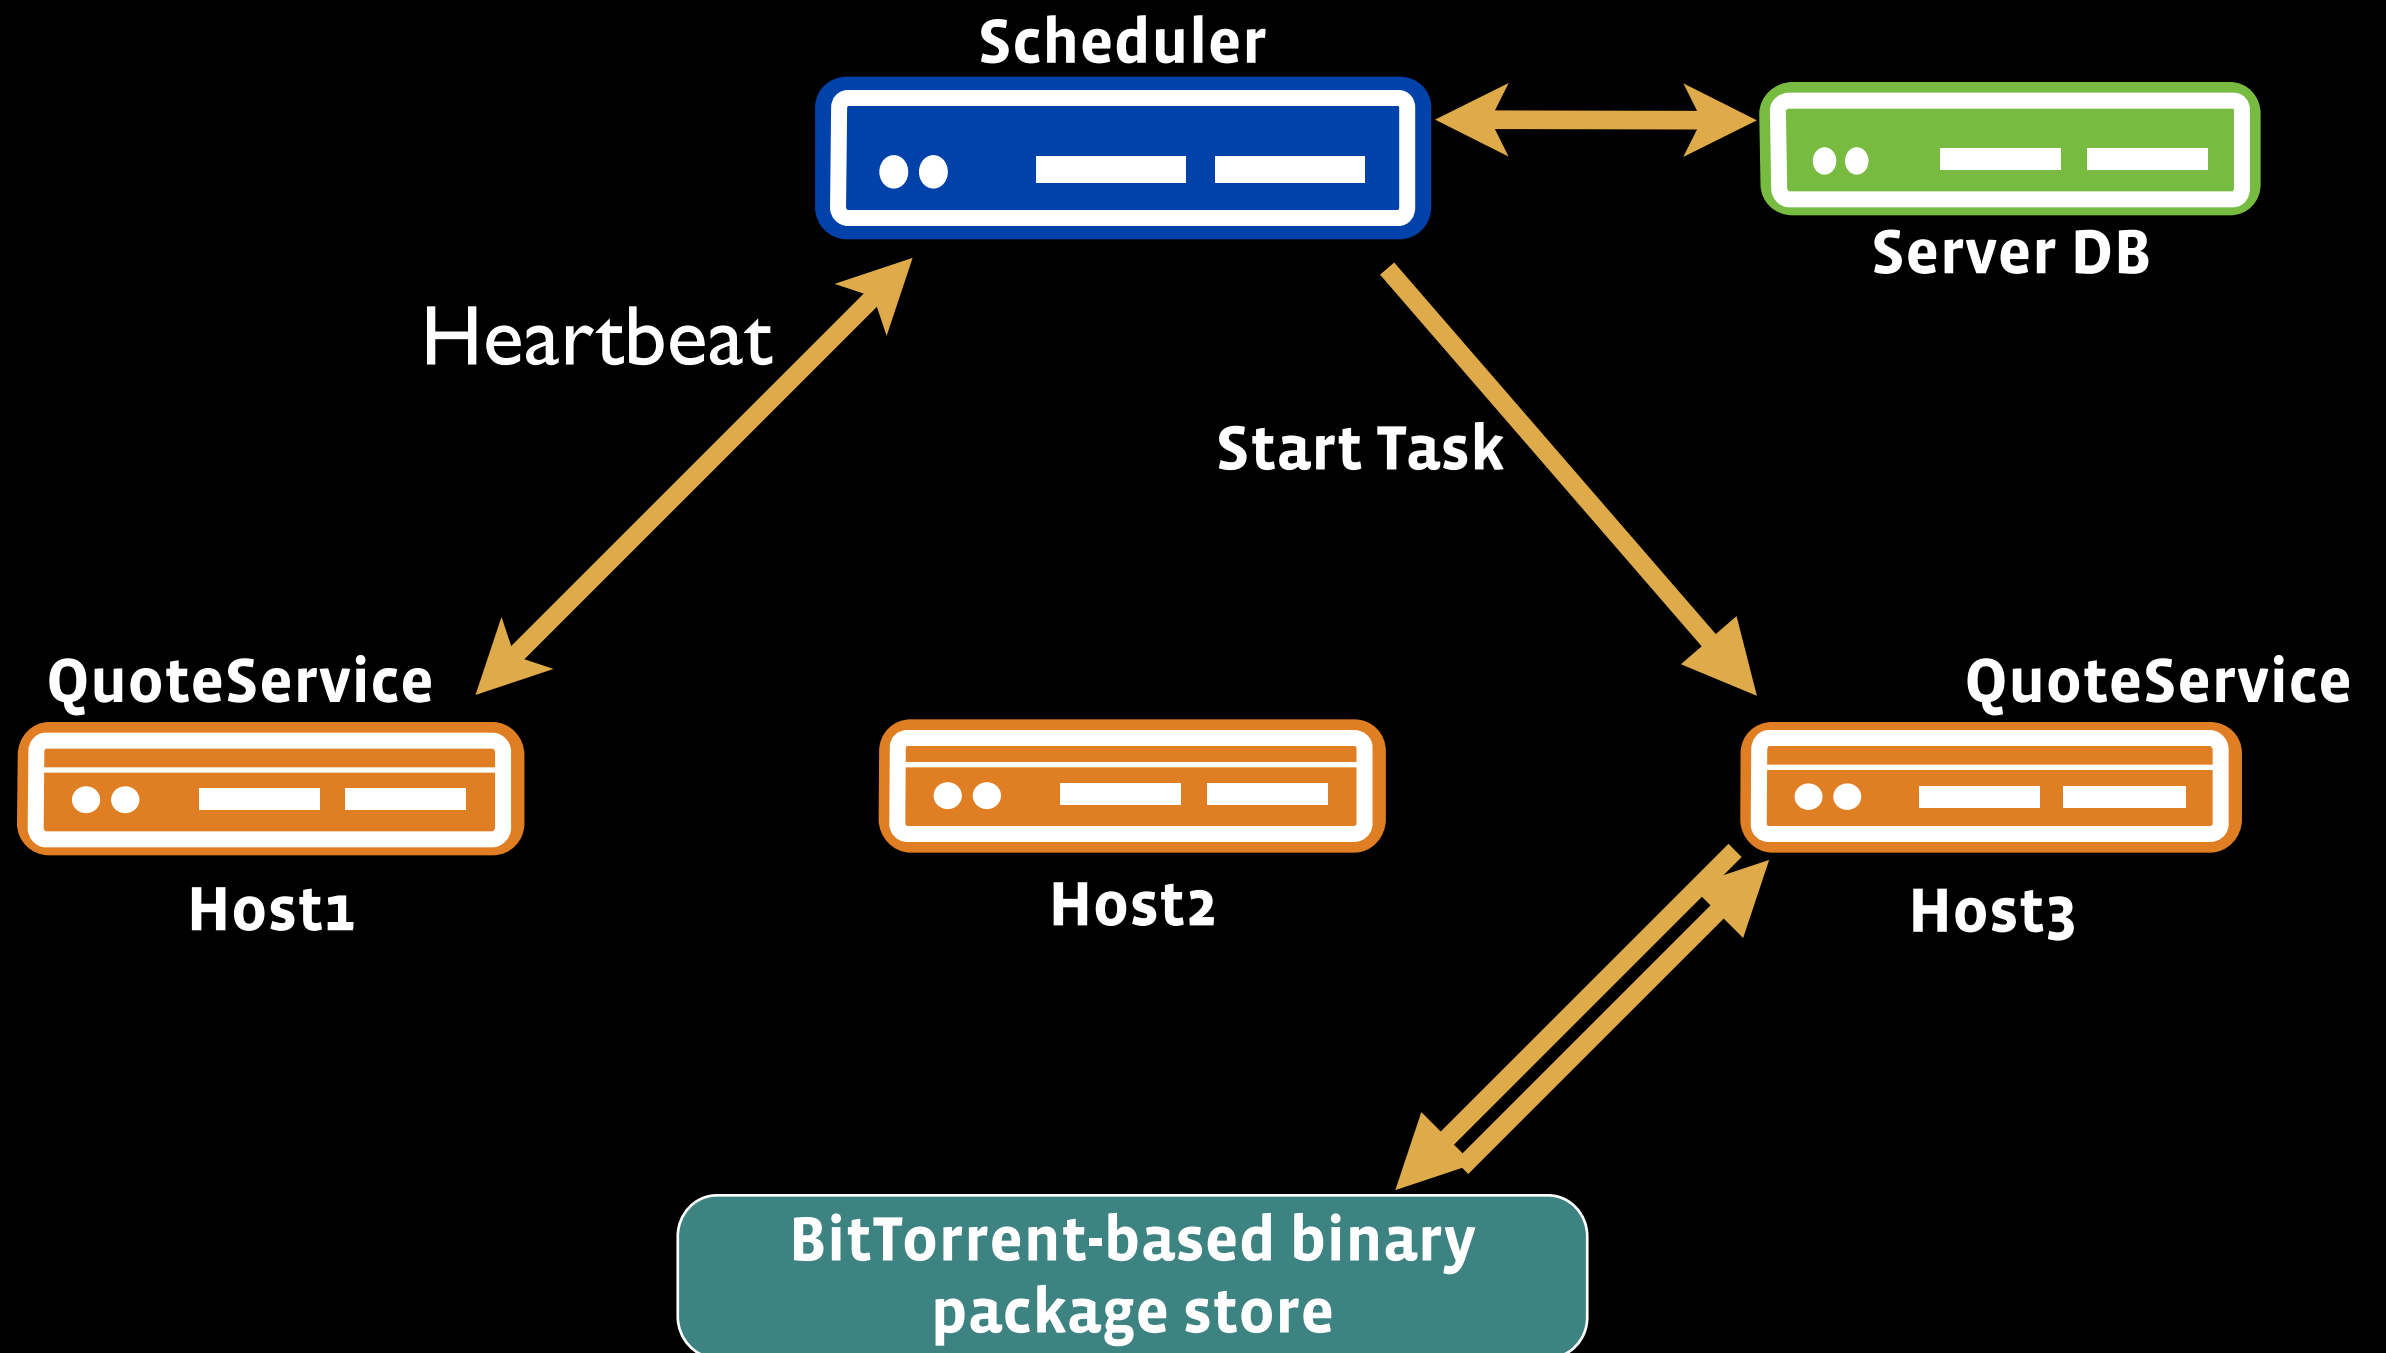


# Failover



Machine “failures” / hour

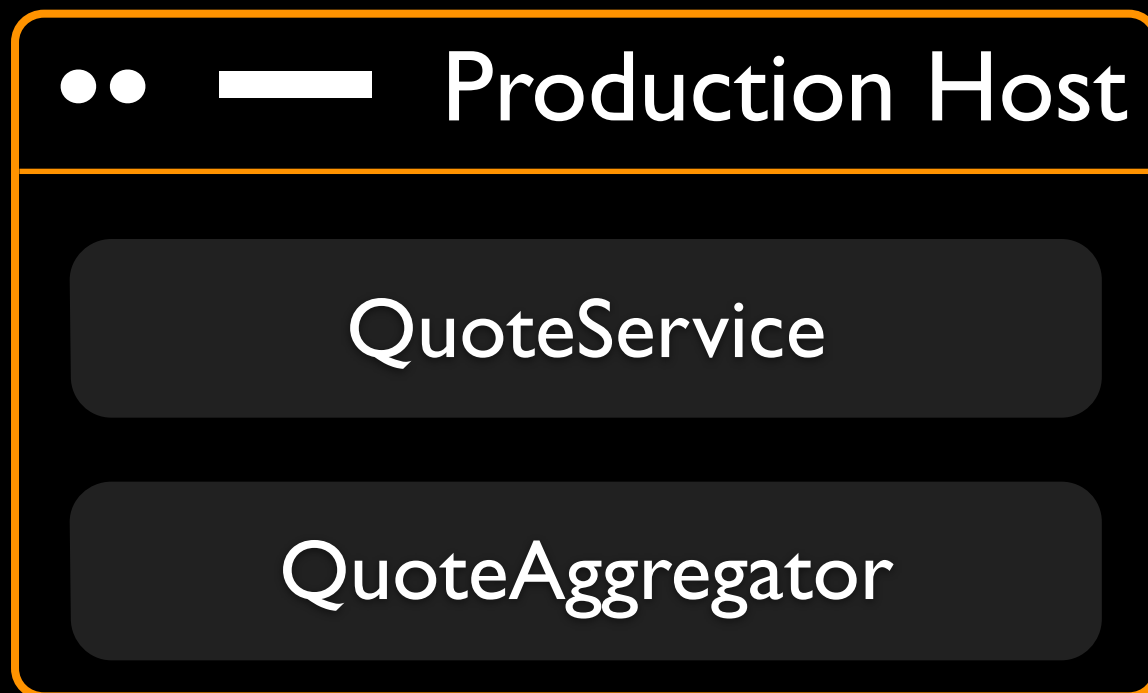

# Failover



# Painless Hardware maintenance

- Notify scheduler of impending operations
- Scheduler can preemptively move tasks
  - Graceful migration for stateless services
  - Stateful services may endure maintenance

# Expressive allocation policies



# TW Agent

•• — — Production Host

TW Agent process

Task Manager

API

Resource Manager

Package Manager

Scheduler heartbeat

Agent Helper

Task A

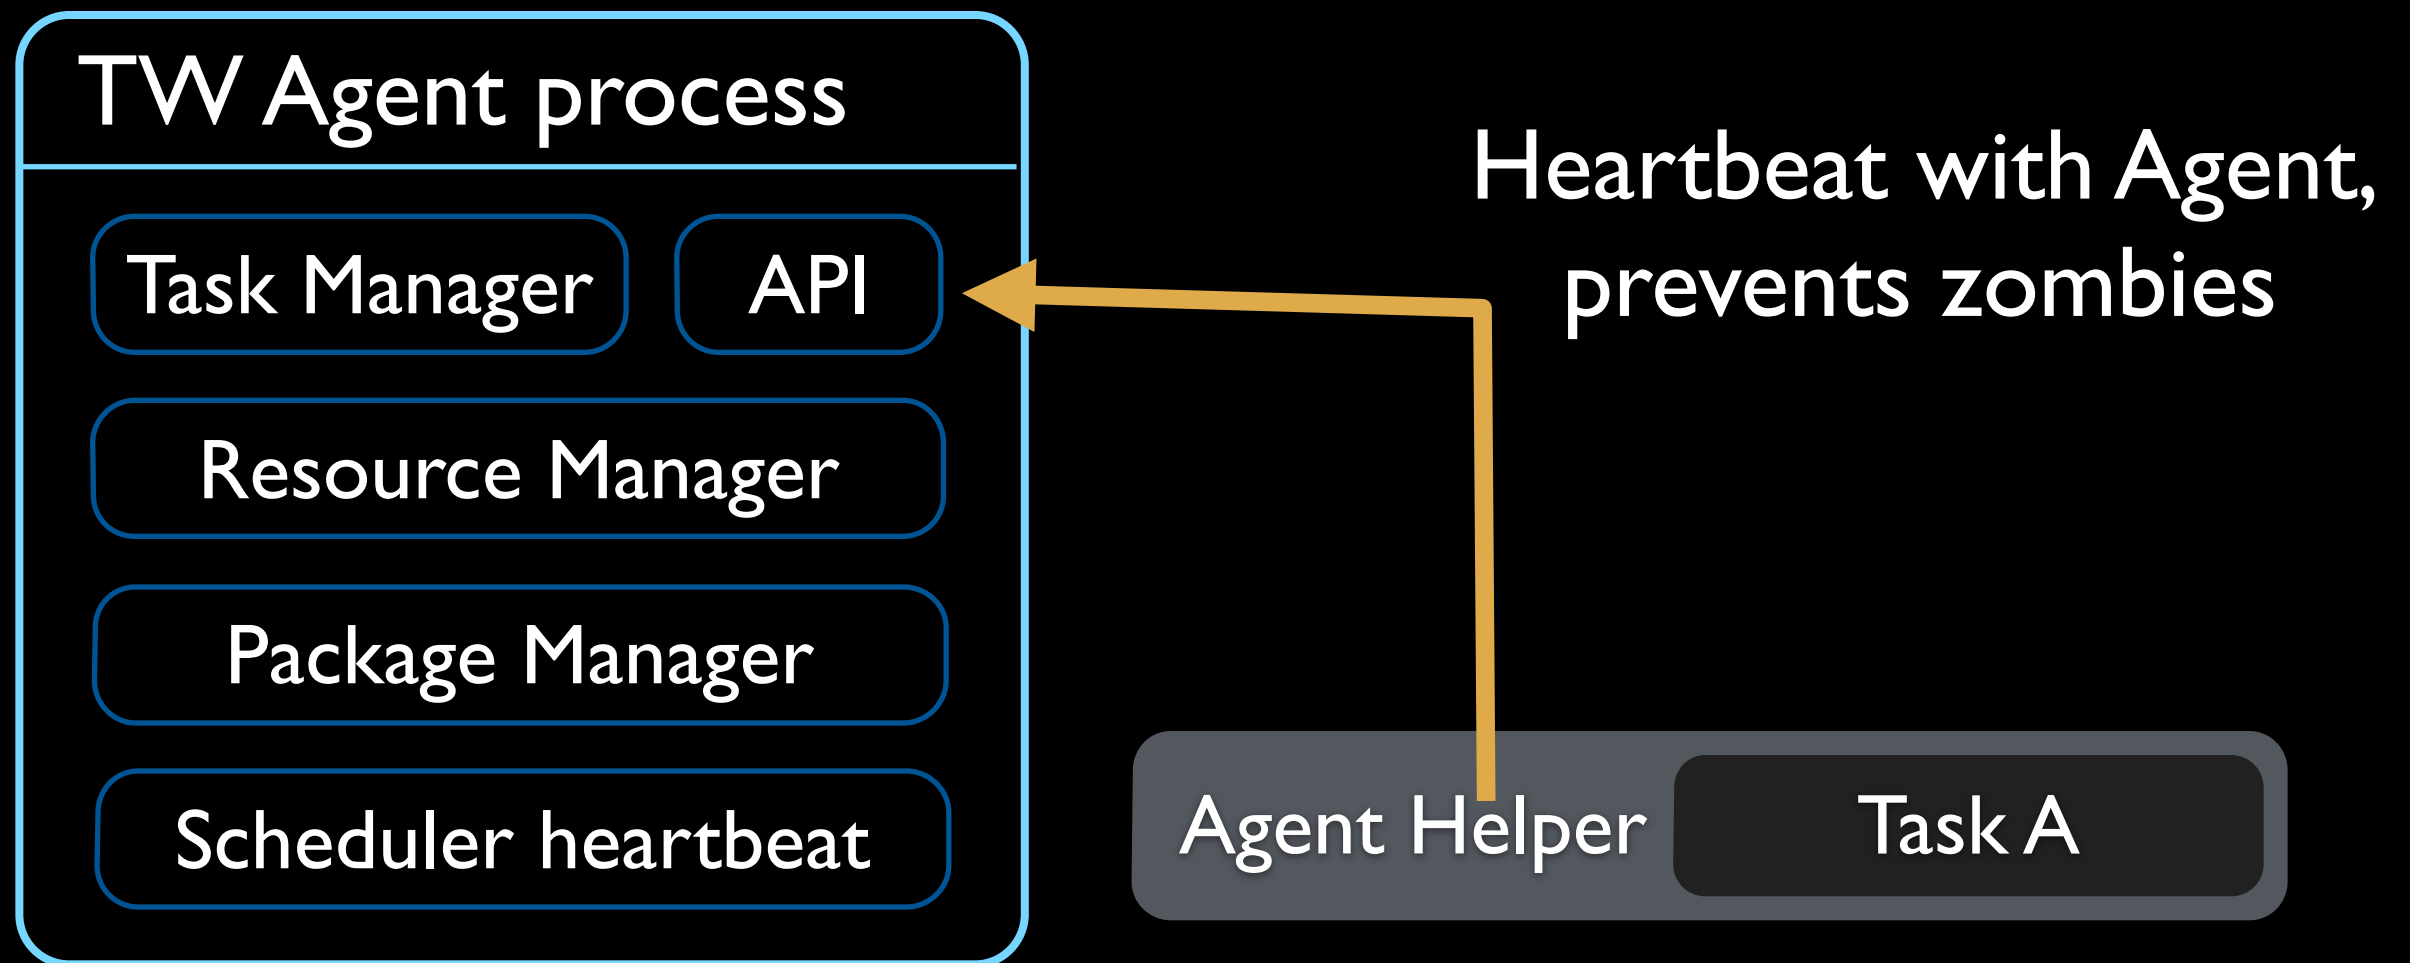
Agent Helper

Task B

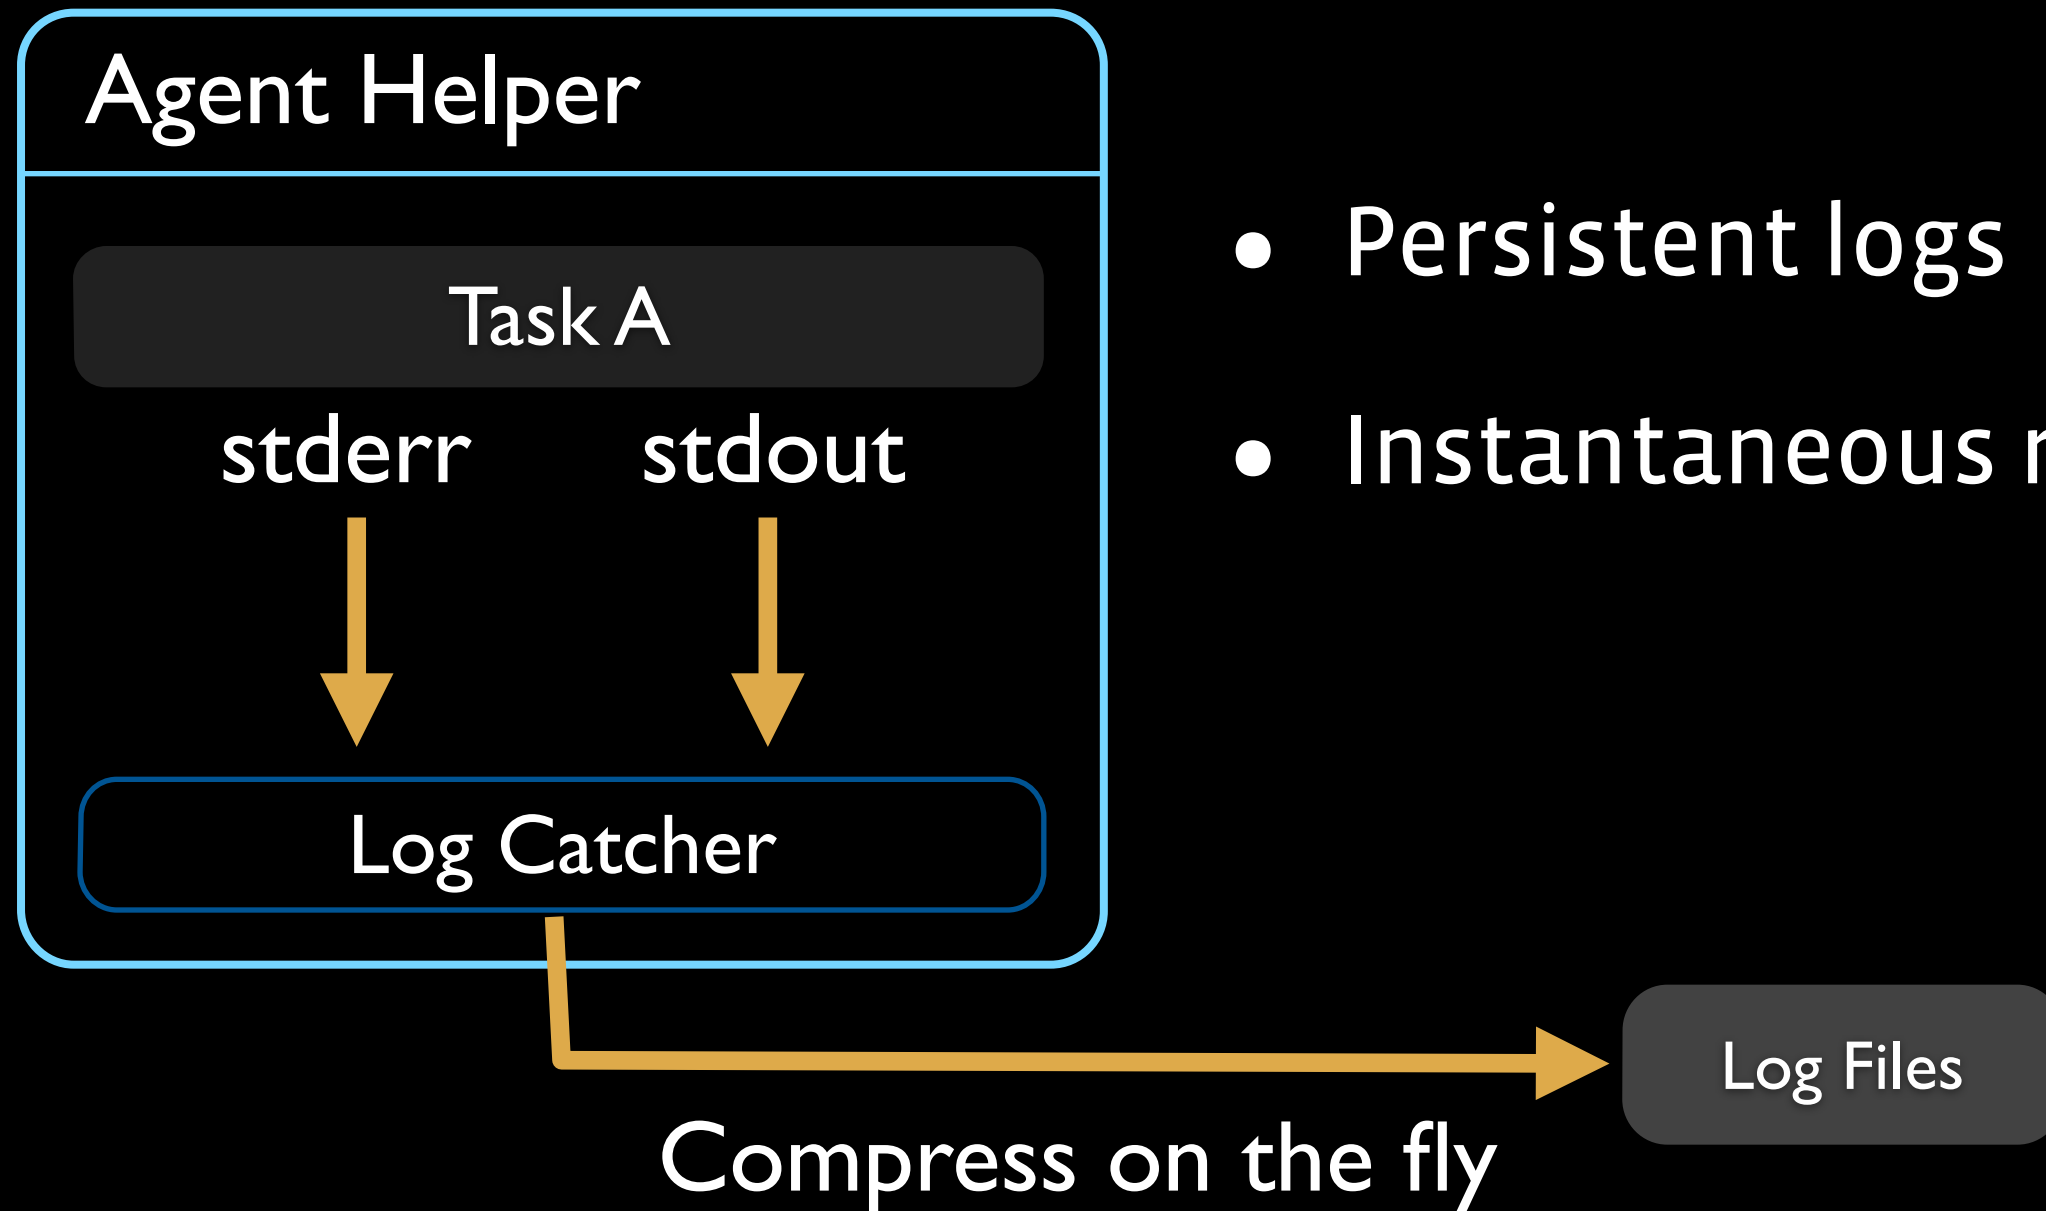
Agent Helper

Task C

# Agent Helper process



# Logging



# Sandboxing

Initially, used `chroots` to contain processes

- No isolation
- Not secure

```
#!/usr/bin/perl -w
chdir "/"; opendir JAILROOT, ".";
mkdir "mysubdir"; chdir "mysubdir";
chroot "."; chdir(*JAILROOT);
while ((stat("."))[0] != (stat(".."))[0] or
       (stat("."))[1] != (stat(".."))[1]) {
    chdir "..";
}
chroot ".";
system("/bin/sh");
```

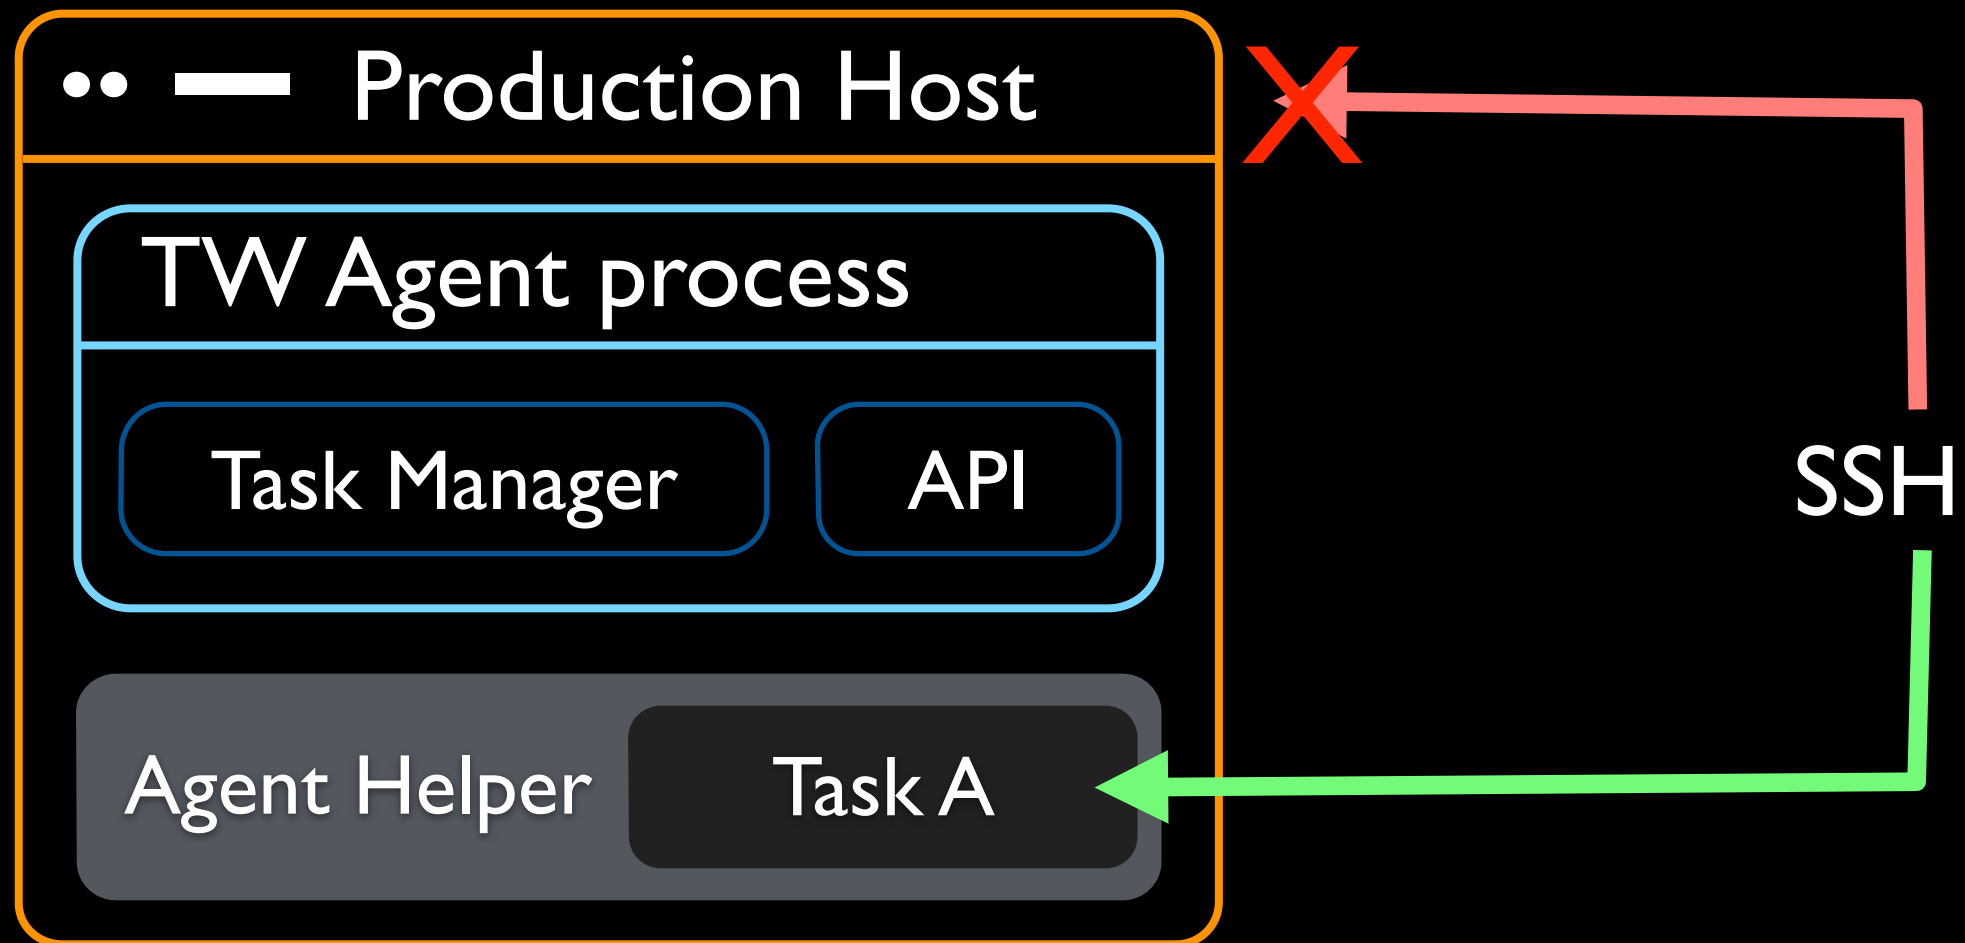


# LinuxContainers

- As tech matured, we switched
- Separate process and file namespaces, set up by Helper
- Mount required resources directly into container
- Secure & isolated

```
syscall(SYS_clone,  
        CLONE_NEWPID | CLONE_NEWUTS |  
        CLONE_NEWIPC | CLONE_NEWNS,  
        0, nullptr, nullptr);
```

# Service permissions



- Every container runs `sshd`
  - SSH directly into the container
  - Regulate access

# Configuring the container

```
import 'git.tw'

job = Job(
  command = '/packages/quote_server/server',
  scheduling = Scheduling(replicas = 10),

  profile = 'centos6', # centos5, ubuntu, ...
  packages = [
    Package(name='quote_server'),
    Package(name='rpm-deps',
             rms=['emacs', 'nfs-utils']),
  ],

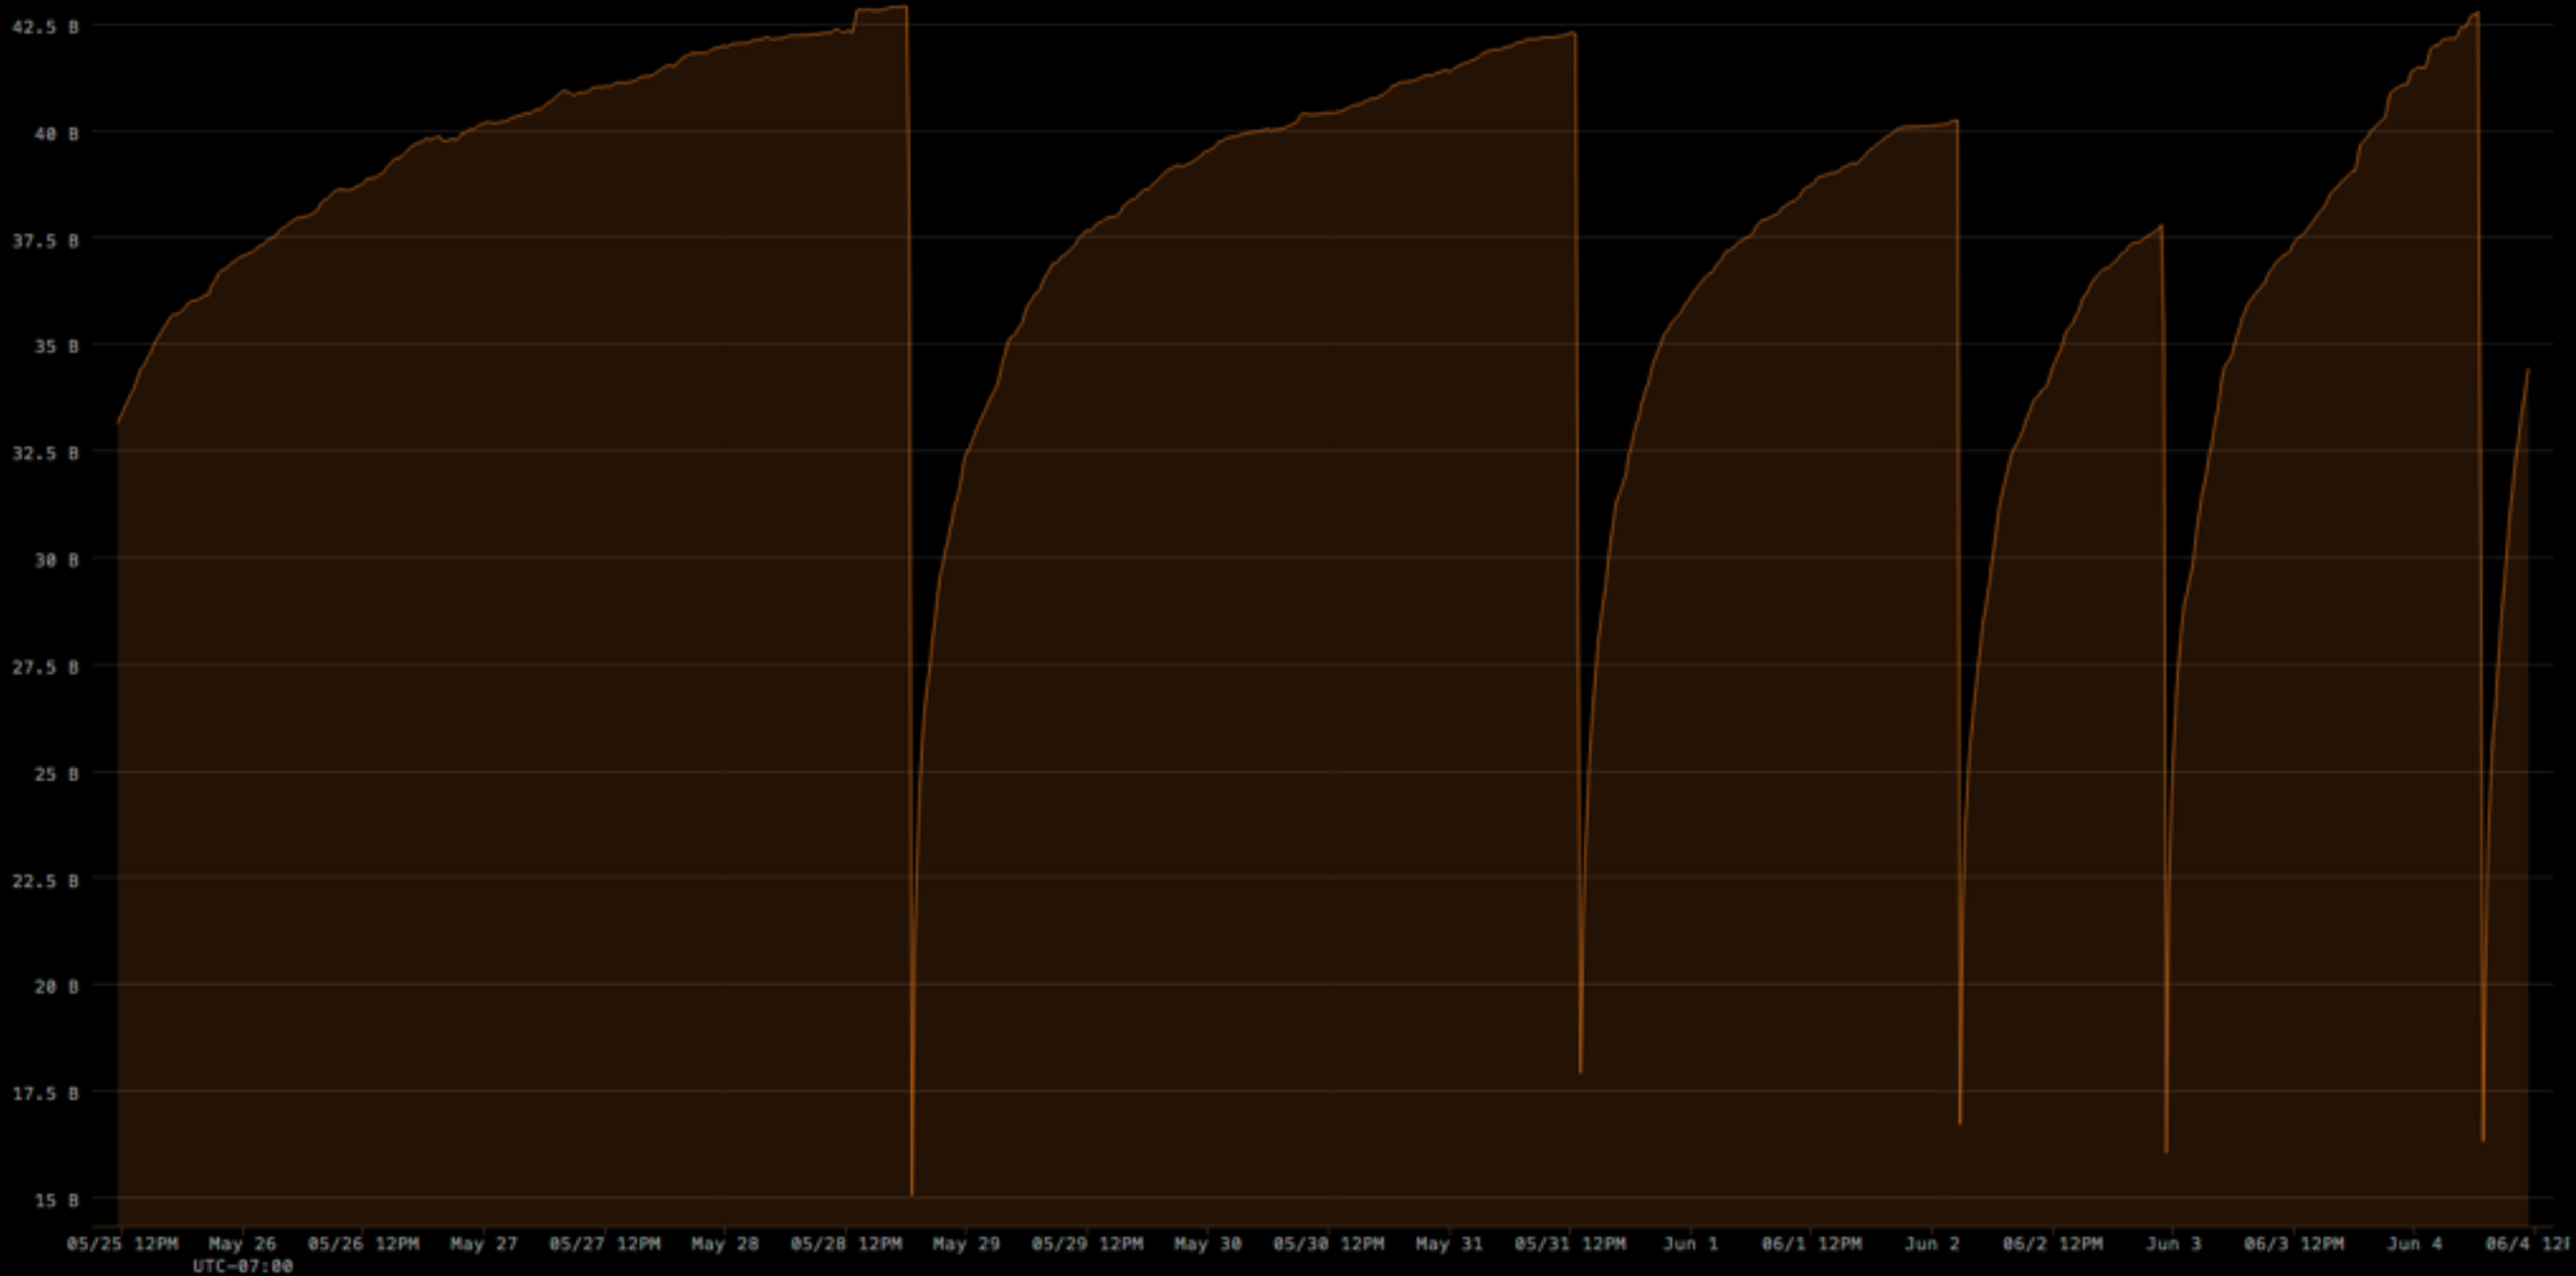
  resource_limits = ResourceLimit(
    cpu = 2, ram = '4G', disk = '32G',
  )
)

enable_git_support(job)
```

# Resource limits

- CPU, RAM & disk limits
- Implemented with `cgroups`
- Agent handles memory limits with `cgroup notification` API
- `Adaptive` limits

# Resource Limits in action

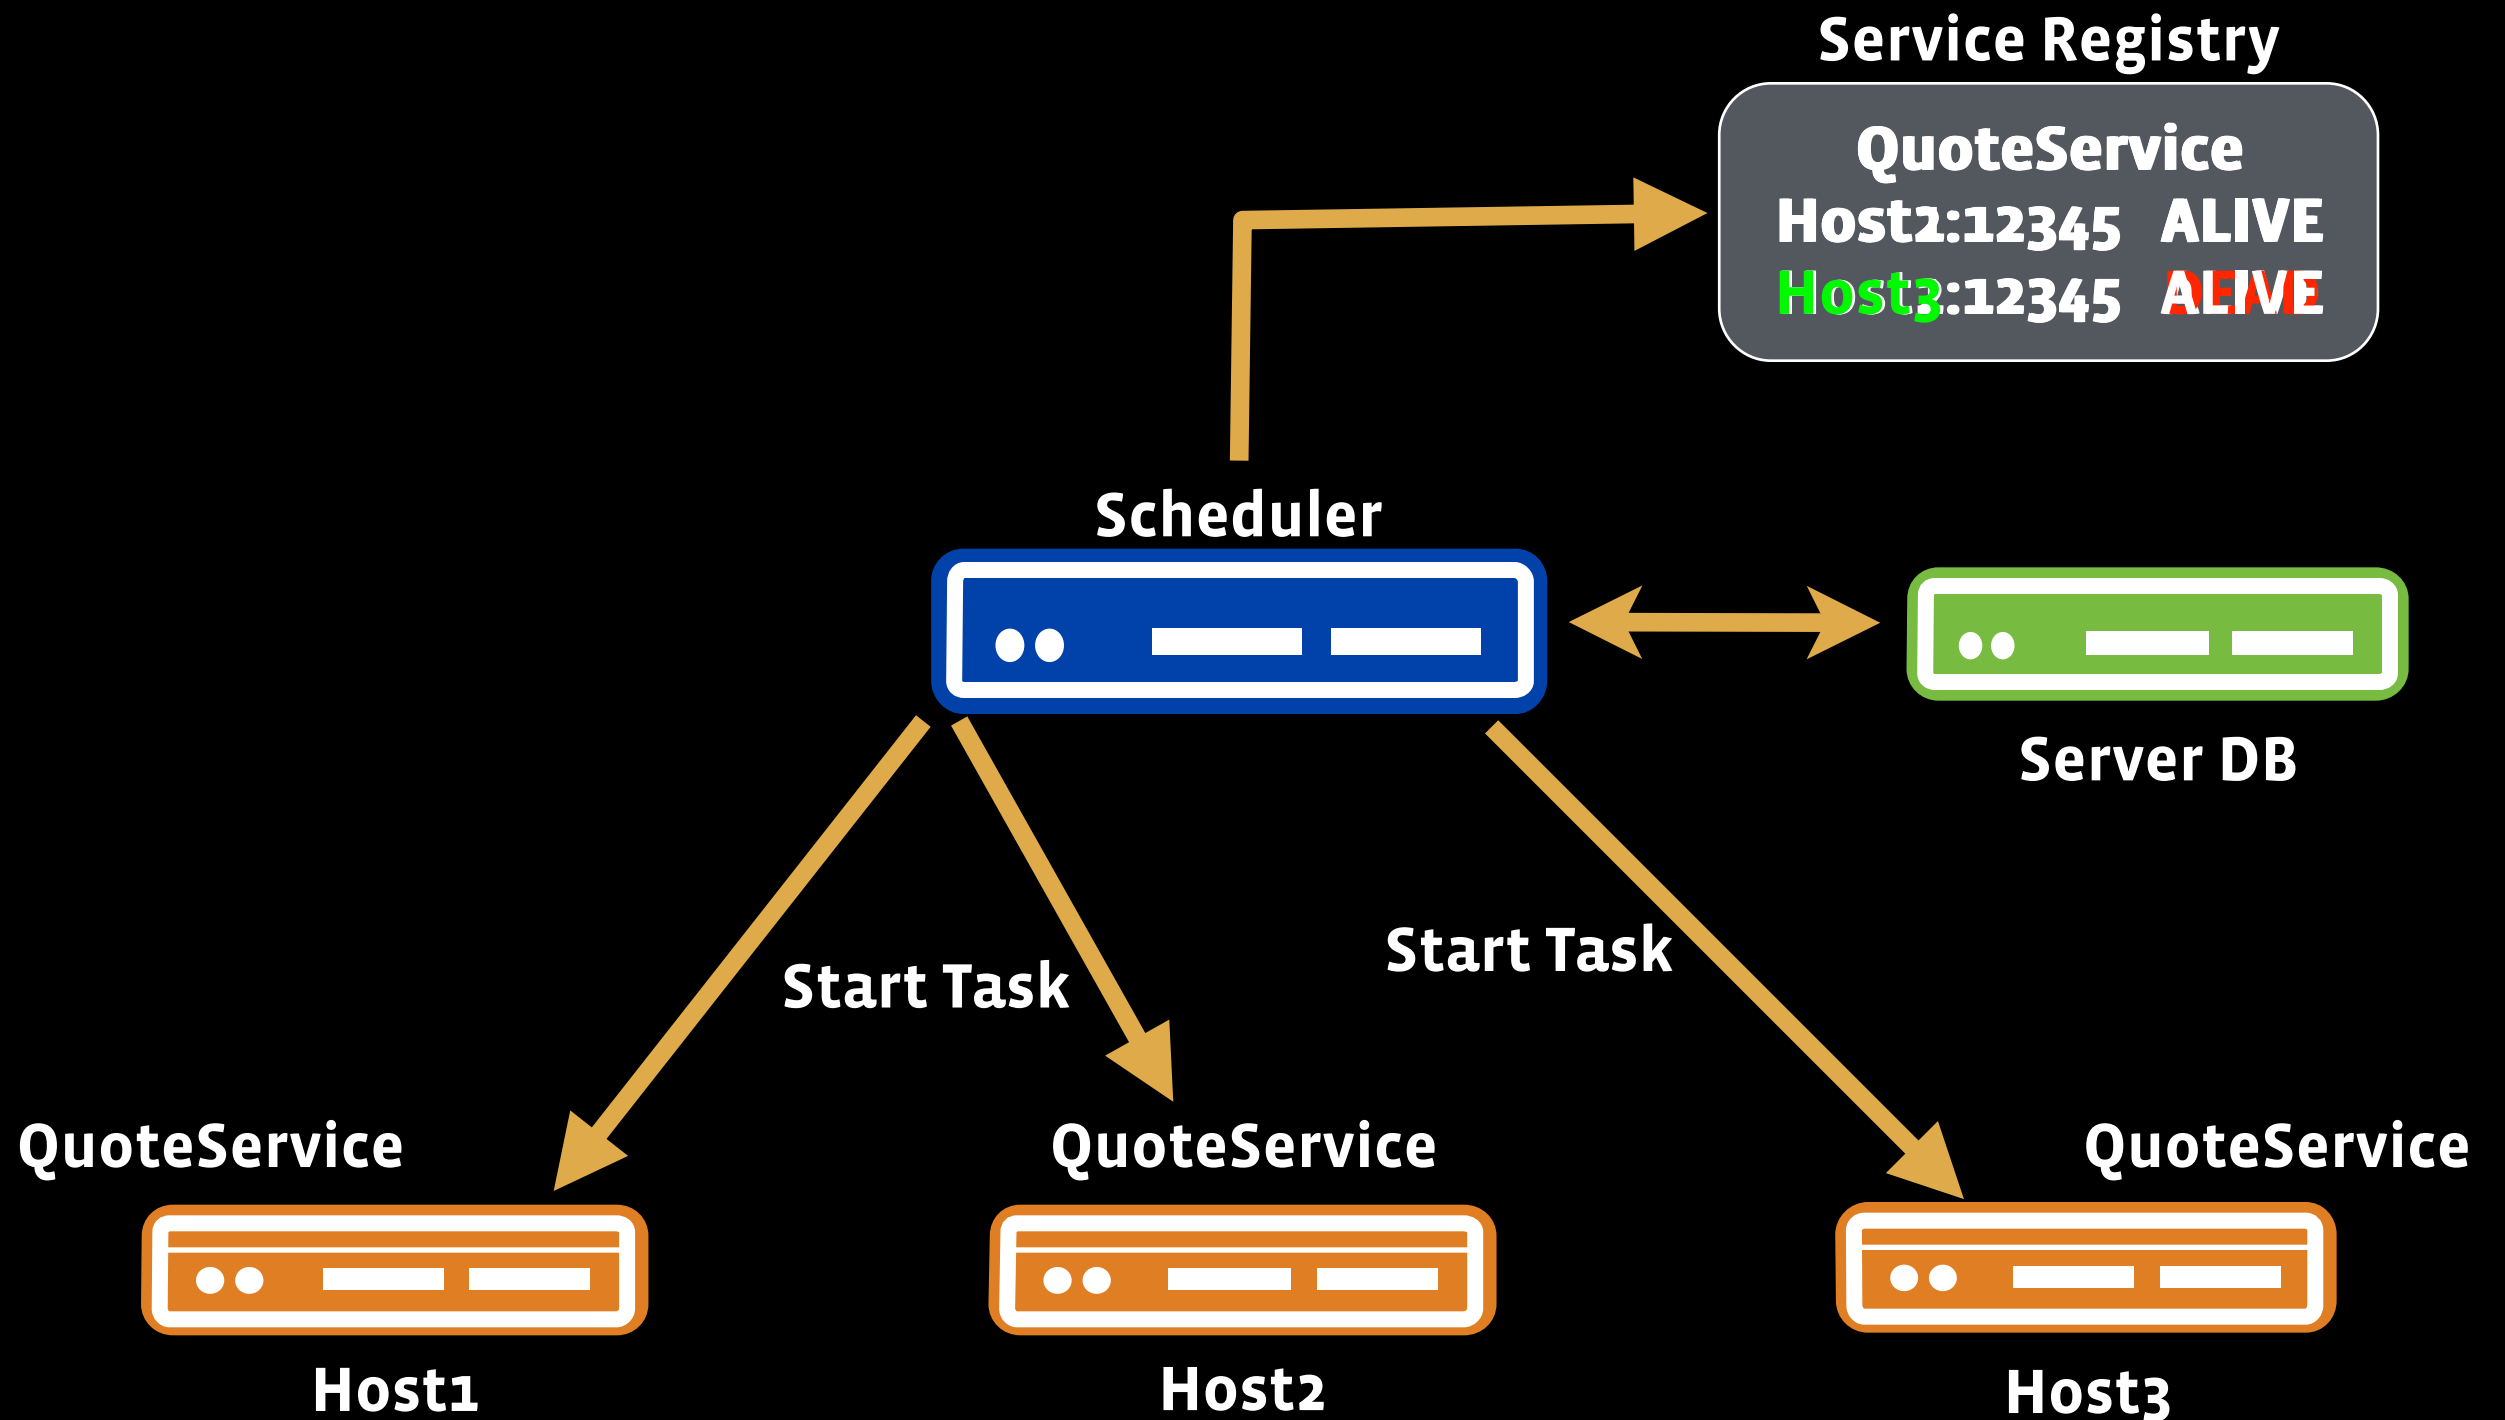


watchdog-service - tw.mem.rss\_bytes

# Migrate from Chroots to Containers

- No-op for most services
- But new namespaces posed problems for some
- Major hurdle was social, not technical

# Service Discovery



# Monitoring & Alerting





# Alternatives to Tupperware

- Why not use Docker / CoreOS?
  - They didn't exist
  - TW integrates with other FB systems
- Why not use VMs?
  - Performance penalty
  - Hypervisor makes debugging harder

Lessons learnt

# Releases are scary!

- Release often
- Dry runs
- Canaries are your friends
- Manage dependencies

# Sane defaults



- Users shouldn't have to read entire manual
- Choose what makes sense for most services

# What went wrong?

- Hard to understand *why* TW did something
- It's not about “what went wrong”, but “what should I do next?”

# Tupperware

- Automated deployment
  - Less work for engineers
- Containers for security and isolation
- Increased efficiency



## Questions?