

DATABASE SHARDING @NETLOG

Tags

performance,
partitioning, federation,
database, php, mySQL,
memcached, sphinx

jurriaanpersyn.com

lead web dev at Netlog

php + mysql + frontend

since 3 years

A Pan European Social Network

NETLOG



Over 40 million active members

Over 50 million unique visitors per month

Over 5 billion page views per month

Over 26 active languages and 30+ countries

6 billion online minutes per month

Top 5 most active countries: Italy, Belgium, Turkey, Switzerland and Germany

Our reach beyond Europe


NETLOG




Visitor statistics

NETLOG

Monthly Visits in EU



Monthly Page Views in EU



- **huge amounts of data** (eg. 100+ million friendships on nl.netlog.com)
- **write-heavy app** (1.4/1 read-write ratio)
- **typical db up to 3000+ queries/sec** (15h-22h)

- most technologies in the web stack are stateless
- the only layer not being stateless is the data itself
- hardest (backend) performance problems were mostly database related

We build Netlog on open source software

NETLOG

- php
- mysql
- apache
- debian
- memcached
- sphinx
- lighttpd
- squid
- and many more ...

Topics

- Netlog history of scaling
- Sharding architecture
- Implications
- Our approach
- Tackling the problems


Master
r/w

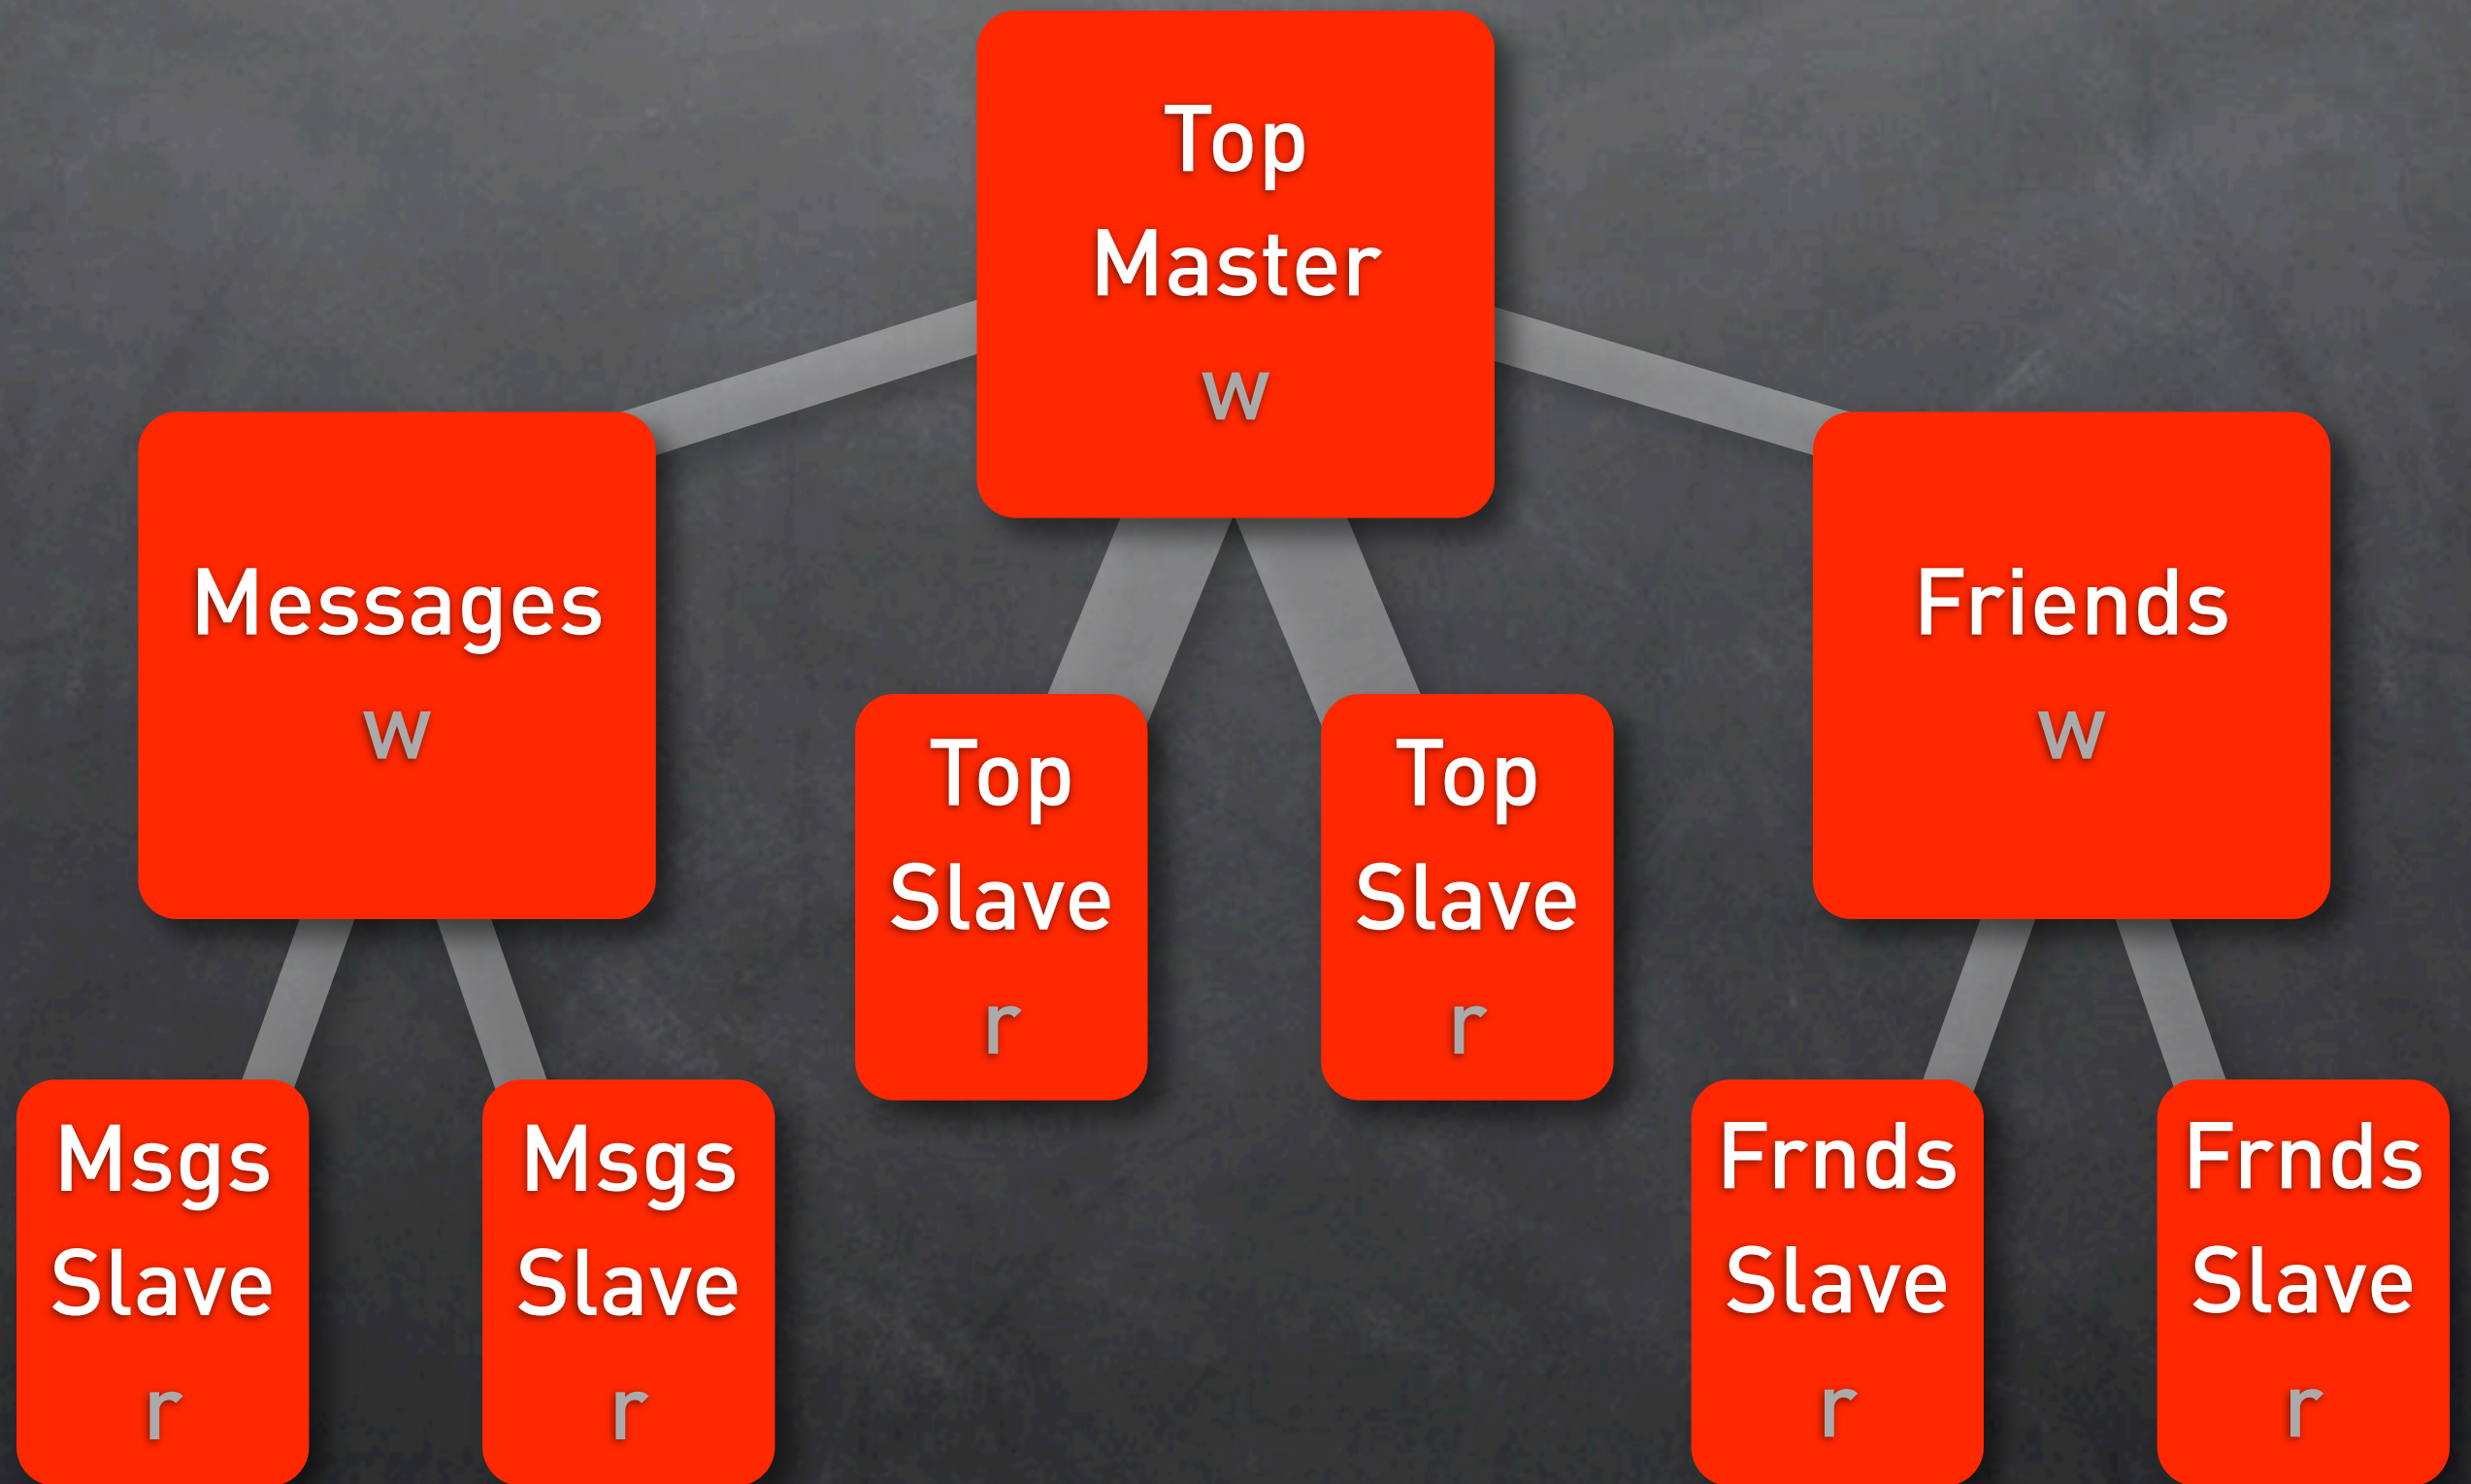

```
graph TD; Master[Master<br/>w] --- Slave1[Slave<br/>r]; Master --- Slave2[Slave<br/>r];
```

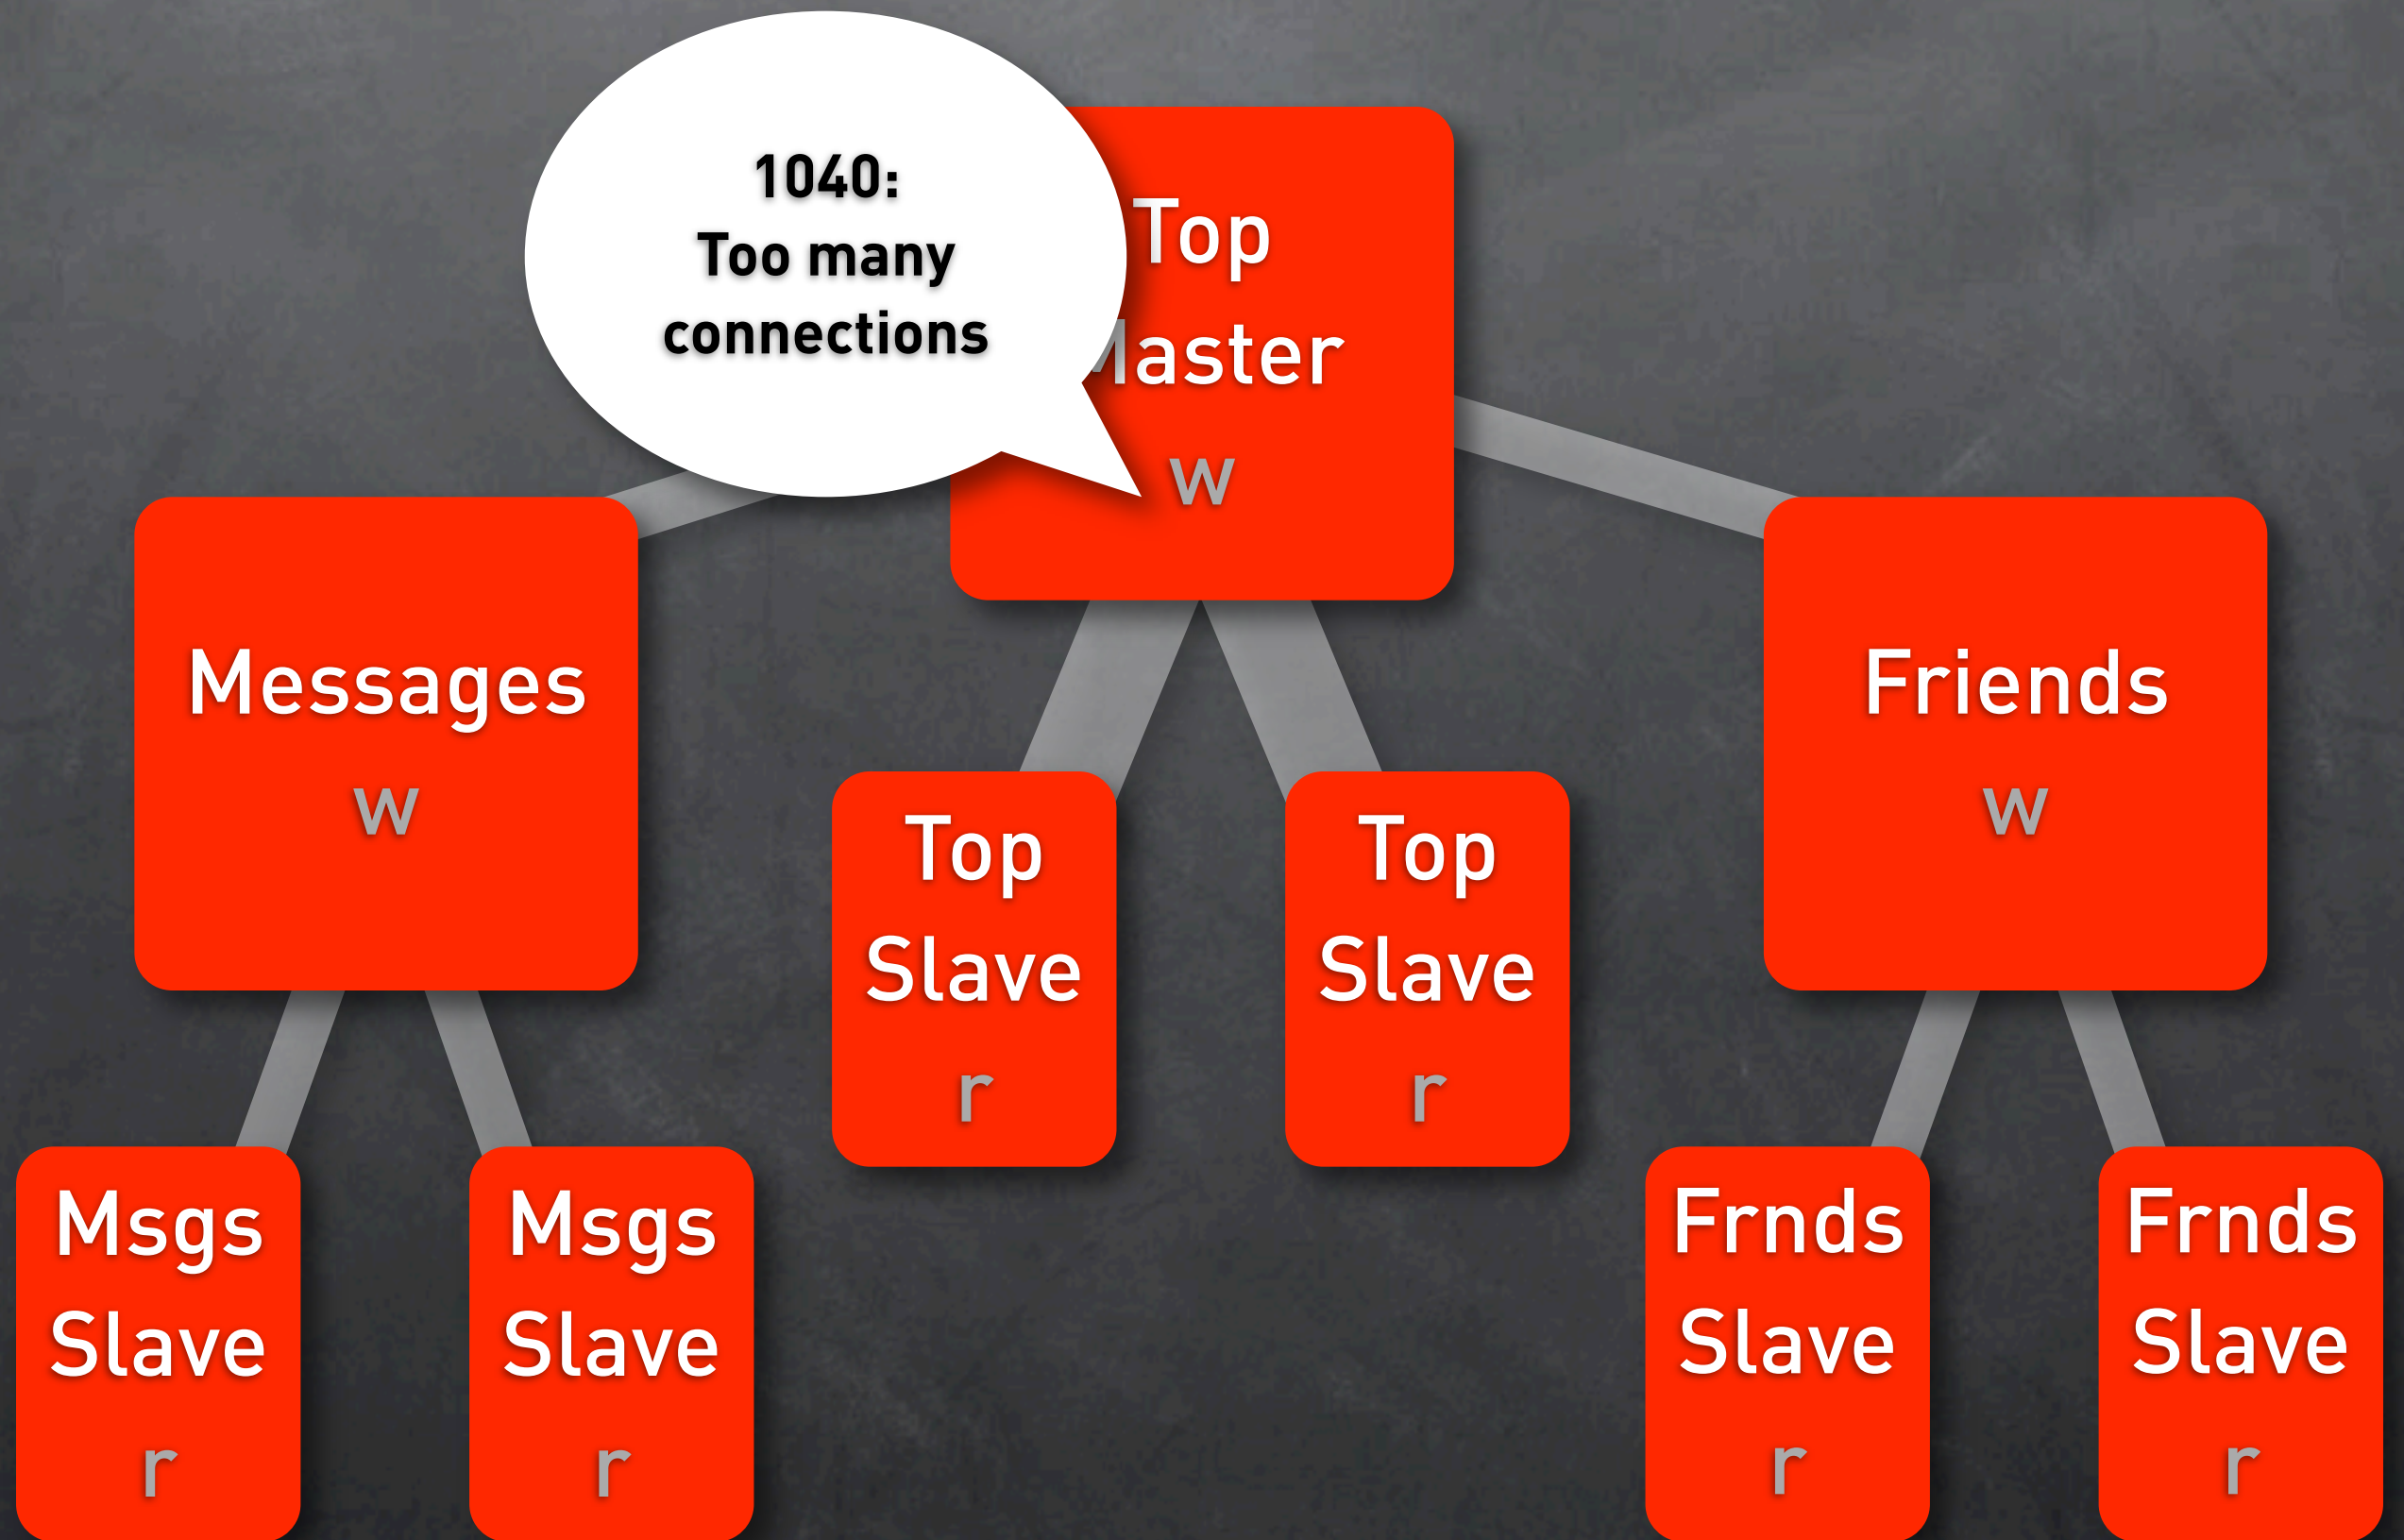
Master
w

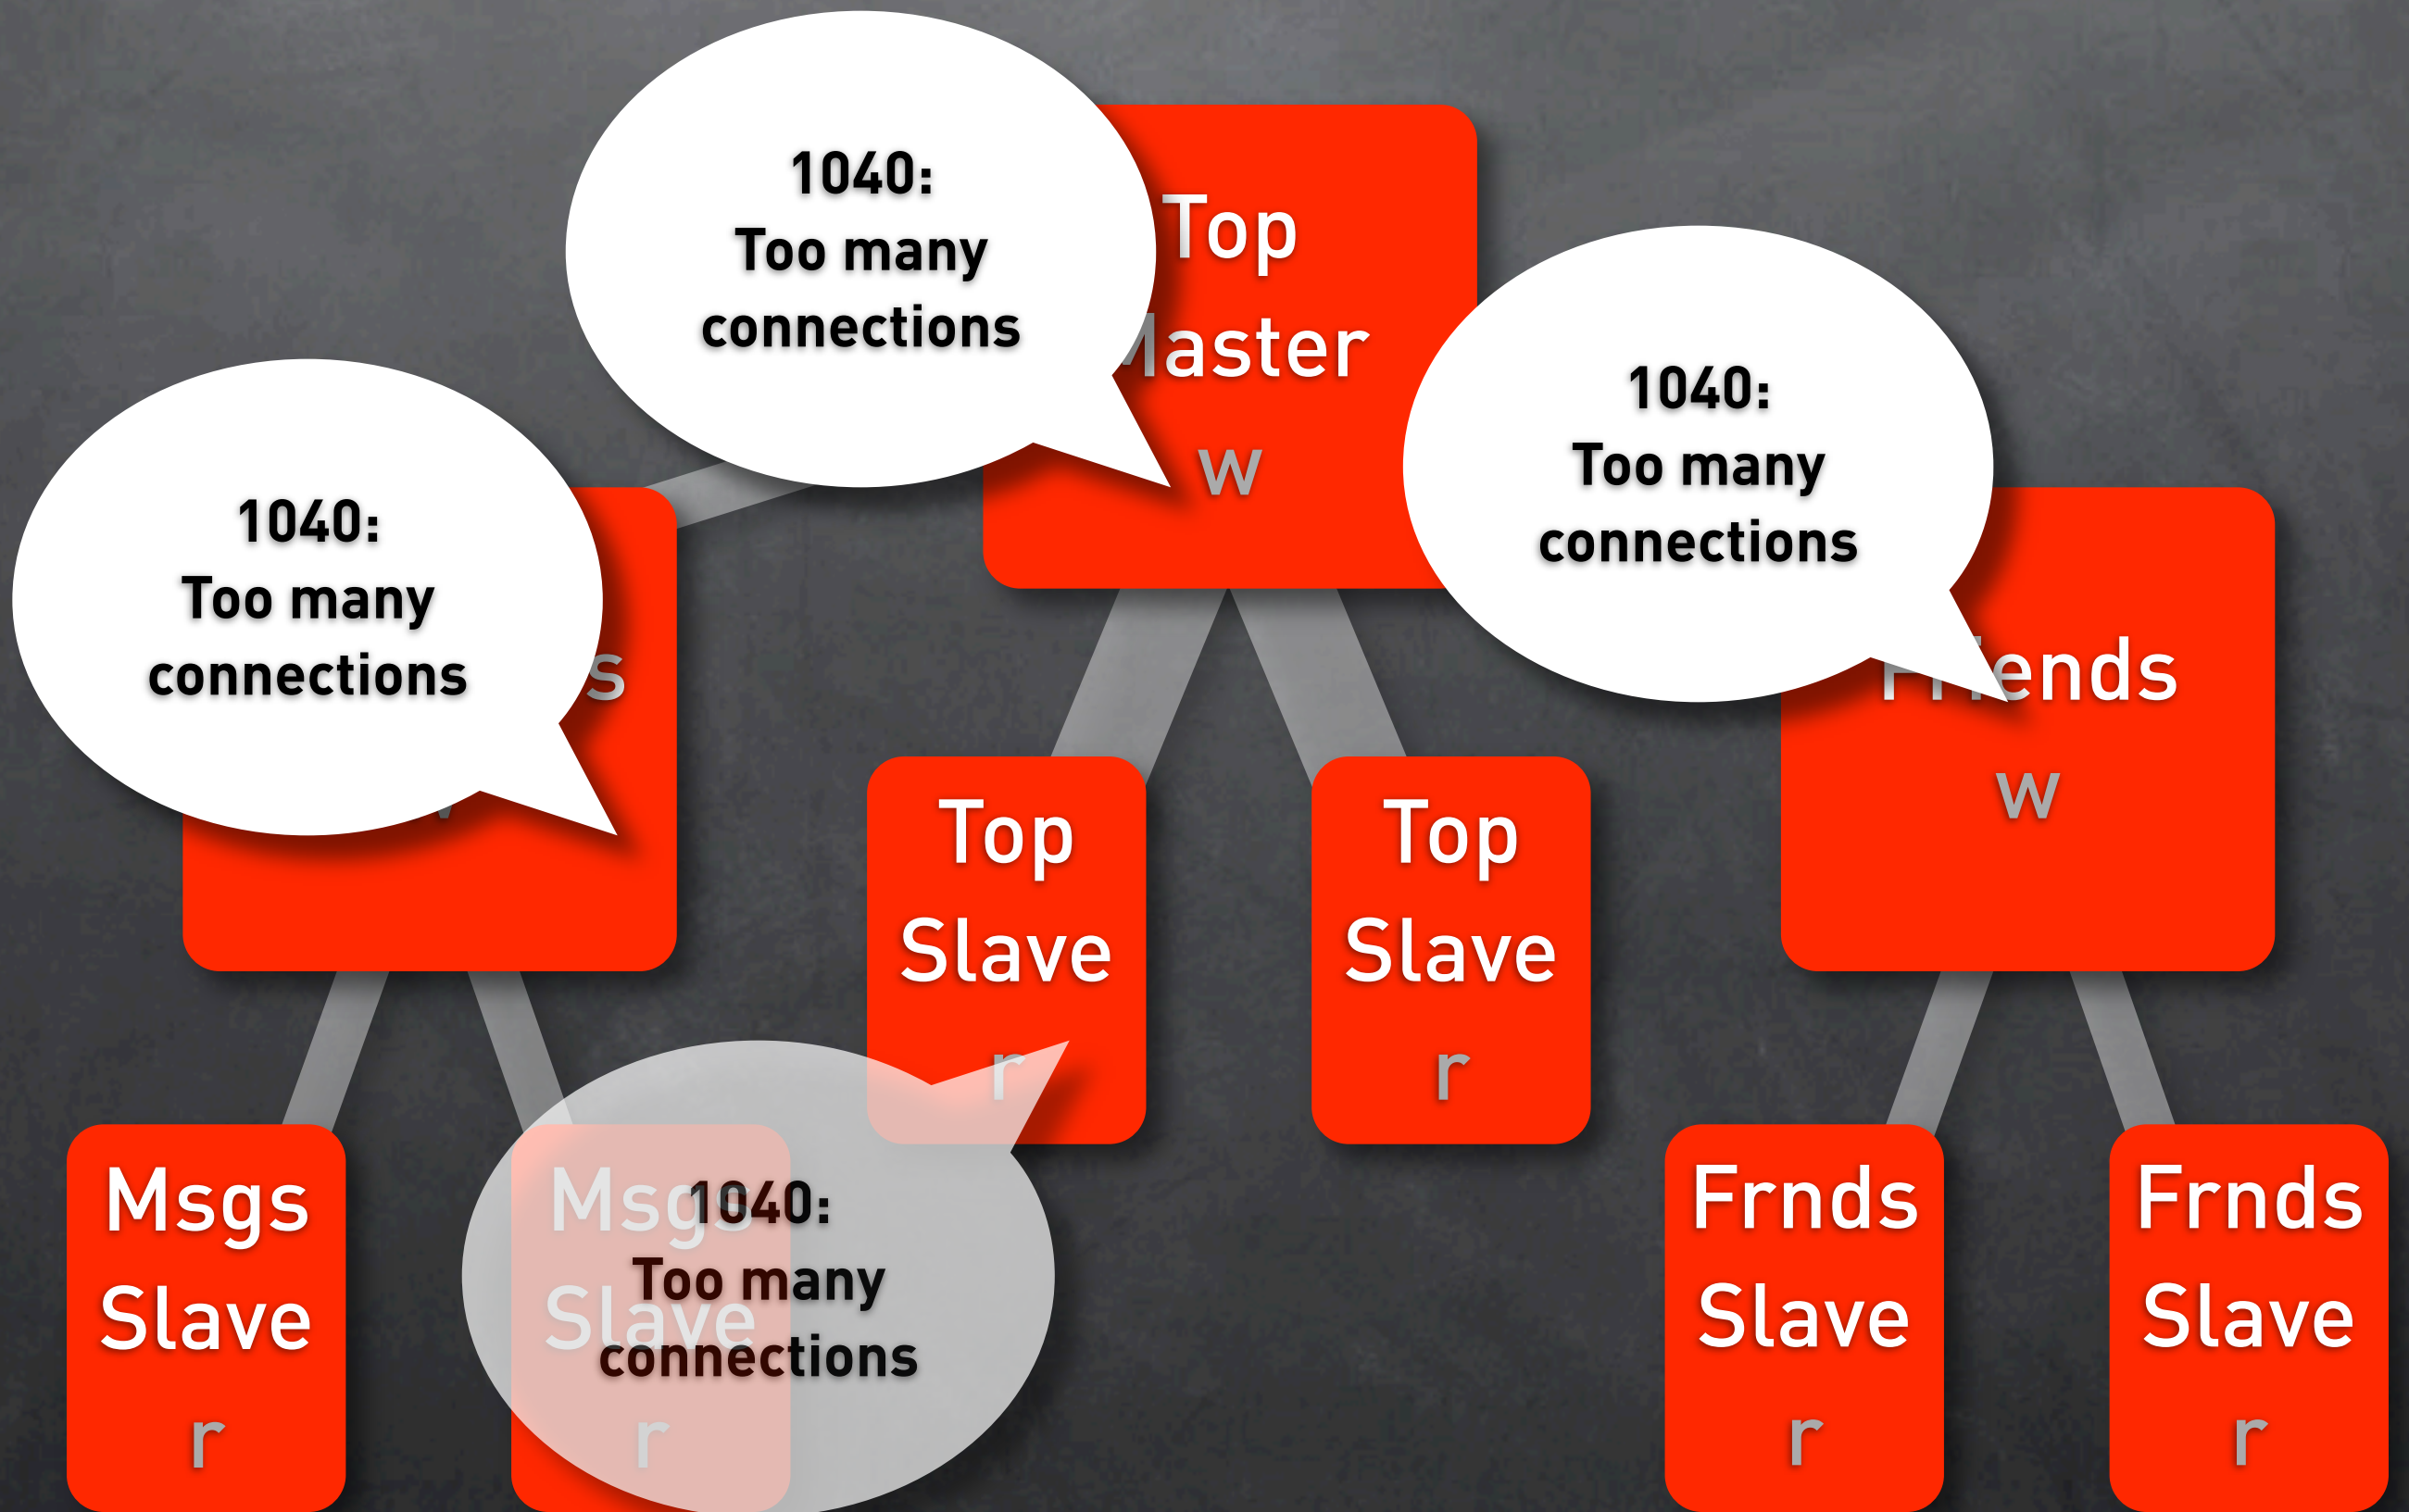
Slave
r

Slave
r














**More vertical
partitioning?**



**Master-to-master
replication?**




Caching?




Sharding!



pieces
fragments
horizontal part.
federation





**More data?
More shards!**



- A blog
- Split users across 2 databases
- Users with even ID go to database 1
- Users with uneven ID go to database 2
- Query: Give me the blog messages from author with id 26


```
$db = DB::getInstance();  
  
$db->prepare("SELECT title, message  
             FROM BLOG_MESSAGES  
             WHERE userid = {userID}");  
  
$db->assignInt('userID', $userID);  
  
$db->execute();
```

```
$db = DB::getInstance($userID);  
  
$db->prepare("SELECT title, message  
              FROM BLOG_MESSAGES  
              WHERE userid = {userID}");  
  
$db->assignInt('userID', $userID);  
  
$db->execute();
```

- “sharding key”
- eg. for a blog
 - partition on publication date?
 - partition on author?
 - author name?
 - author id?
 - partition on category?

- Partitioning schemes
- eg.
 - Vertical
 - Range based
 - Key or hash based
 - Directory based

Several smaller tables means:

- Your users love you again
 - You're actually online!
- Your DBA loves you again
 - Less DB load/machine, less crashing machines
 - Even maintenance queries go faster again!

CAUTION

AREA PATROLLED



By

Ken Danton

Attack Bunny

Security Co.

- Problem: No cross-shard SQL queries
 - `(LEFT) JOIN` between shards becomes impossibly complicated
- Solutions:
 - It's possible to design (parts of) the application so there's no need for cross-shard queries

- Solutions:
 - Parallel querying of shards
 - Double systems
 - guestbook messages:
 - shard on id of guestbook owner
 - shard on id of writer
 - Denormalizing data

- Data consistency / referential integrity?
 - enforce integrity on application level
 - “cross server transactions”
 1. start transactions on both servers
 2. commit transactions on both servers
- check/fix routines can become substantial part of development cost

- Balancing (eg. balanced on a user's ID)
 - What about “power users”?
 - Differences in hardware performance?
 - Adding servers if you grow?
 - Partitioning scheme choice is important
 - directory based++, but: SP0F? overhead?

- Your web servers talk to more / several databases
- network implications
- keep connections open? close after every query? pools?

- **MySQL Cluster** (high availability, performance, not distribution of writes)
- **MySQL Partitioning** (not feature complete when we needed it, Netlog not 5.1 ready/compatible at that time, not directory based (??))
- **HiveDB** (mySQL sharding framework in Java, requires JVM, php interface in infancy state)

- **MySQL Proxy** (used by following two options, introduces LUA)
- **HSCALE** (not feature complete, partitions files, not yet cross-server, builds on MySQL Proxy)
- **Spock Proxy** (fork of MySQL Proxy, atm only range based partitioning)
- **HyperTable** (HQL)
- **HBase / BigTable**

- **Hibernate Shards**
(whole new DB layer for us)
- **SQLAlchemy**
(for Python)
- **memcached from Mysql**
(SQL-functions or storage engine)
- **Oracle RAC**
(well, not MySQL)

- flexible for your hardware department
- no massive rewrite
- incremental implementation
- support for multiple sharding schemes
- easy to understand
- well balanced

- in-house
- 100% php
- middleware between application logic and `class` DB
- complete caching layer built in
- most sharded data carved by `$userID`

sharddbhost001

sharddb001

shard0001

shard0002

shard0003

shard0004

sharddb002

shard0005

shard0006

shard0007

shard0008

- **Sharding Management**
 - **Lookup System** (directory based)
 - **Balancer / Manager**
- **Sharded Tables API**
 - **Database Access Layer**
 - **Caching Layer**

- Lookup system translates `$userID` to the right db connection details
- `$userID` to `$shardID`
(via SQL/memcache - combination not fixed!)
- `$shardID` to `$hostname` & `$databasename`
(generated configuration files, with flags about shard being available for read/write)

- All sharded records have a
 - “shard key” (\$userID)
 - “\$itemID” (identification of the record)
 - related to `$userID`
 - mostly: combined auto_increment column

- Example API:
 - An object per `$tableName/$userID`-combination
 - implementation of a class providing basic CRUD functions
 - typically a class for accessing database records with “a user’s items”

Query: Give me the blog messages from author with id 26. 3 queries:

1. where is user 26? › shard 5
2. on shard 5 › give me all “blogIDs” (itemIDs) of user 26 › blogID 10, 12 & 30
3. on shard 5 › fetch details WHERE blogID IN(10,12,30) › array of title + message

- Define 'load' percentage for shards (#users), databases (#users, #filesize), hosts (#sql reads, #sql writes, #cpu load, #users, ...)
- Balance loads and start move operations
 - We can do this on userID level
 - completely in PHP / transparant / no user downtime

- Downtime of a single databasehost affects only users on shards on that DB
- a few profiles not available
- communication with that profile not available

- memcached
- parallel processing
- sphinx

- Makes it faster
 - much faster
 - much much faster
- Makes some “cross shard” things possible

```
function isObamaPresident()
{
    $memcache = new Memcache();
    $result = $memcache->get('isobamapresident'); // fetch
    if ($result === false)
    {
        // do some database heavy stuff
        $db = DB::getInstance();
        $votes = $db->prepare("SELECT COUNT(*) FROM VOTES WHERE
vote = 'OBAMA'")->execute();
        $result = ($votes > (USA_CITIZEN_COUNT / 2)) ? 'Sure
is!' : 'Nope.'; // well, ideally
        $memcache->set('isobamapresident', $result, 0);
    }
    return $result;
}
```

- “Shard key” to “shard id” is cached
- Each sharded record is cached as array
(key: table/userID/itemID)
- Caches with lists, and caches with counts
(key: where/order/...-clauses)
- “Revision number” per table/shard key-combination

Query: Give me the blog messages from author with id 26. 3 memcache requests:

1. where is user 26? › +/- 100% cache hit ratio
2. on shard 5 › give me all “blogIDs” (itemIDs) of user 26 › list query (cache result of query)
3. on shard 5 › fetch details WHERE blogID IN(10,12,30) › details of each item +/- 100% cache hit ratio (multi get on memcache)

- What? Cached version number to use in other cache-keys
- Why? Caching of counts / lists
- Example: cache key for list of users latest photos (simplified): `"USER_PHOTOS" . $userID . $cacheRevisionNumber . "ORDERBYDATEADDDDESCLIMIT10"`;
- `$cacheRevisionNumber` is number, bumped on every CUD-action, clears caches of all counts +lists, else unlimited ttl.
- "number" is current/cached timestamp

- Several caching modes:
 - `READ_INSERT_MODE`
 - `READ_UPDATE_INSERT_MODE`
- More PHP processing
 - Needs memory
 - PHP-webservers scale more easily

- Split single HTTP-request into several requests and batch process (eg. Friends Of Friends feature)
- Fetching friends of friends requires looping over friends (shard) (memcache makes this possible)
- But: people with 1000+ friends › Process in batches of 500?
- Processing 1000+ takes longer then 2*500+

- **Problem:**
How do you give an overview of eg. latest photos from different users? (on different shards)
- **Solution:**
Check Jayme's presentation "Sphinx search optimization", distributed full text search.
(Use it for more than searching!)

- First and foremost
 - Don't do it, if you don't need to!
(37signals.com)
 - Shard early and often!
(startuplessonslearned.blogspot.com)

- “Don’t do it, if you don’t need to!”?
- Sharding isn’t that easy
- There is no out of the box solution that works for every set-up/technology/...
- Maintenance does get a bit tougher
- Complicates your set-up
- There might be cheaper and better hardware available tomorrow.

- “Shard early and often!”?
- What is the most relevant key to shard on for your app? (eg. userID)
- Do you know that key on every query you send to a database? (function calls, objects)
- Design your application / database schemas so you know that key everywhere you need it. (Migrating schemas is also hard.)

- Don't forget server tuning / hardware upgrade / sql optimization
 - can be easier than (re-)sharding
- Only scale those parts of your application that require high performance
- (some clutter can remain on your db's, but are they causing harm?)

- Migration to sharding system
 - Each shard reads from the master and copies data (can run in parallel, eg. from different slaves, to minimize downtime)
 - Usage of Sharded Tables API is possible without tables actually been sharded (transition phase).

- We don't need super high end hardware for our database systems anymore. Saves \$.
- Backups of your data will be different.
- You can run alters/backups/maintenance queries on (temp) slaves of a shard and switch master/slave to minimize downtime.
- If read-write performance isn't an issue to shard your data, maybe the downtime it takes for an ALTER query on a big table can be.

- We didn't have to balance user by user yet
 - power users balanced inactive users
 - “virtual shards” allowed us to split eg. 1 host with 12 db's into 2 hosts with 6 db's
- Have a plan for scaling up.
 - When average load reaches x how will you add?
- Goes together w/ replication+clustering

**don't do it if you don't need to,
but prepare for when you do**

**the last thing you want is
to be unreachable due to popularity**

netlog.com/go/developer

jurriaan@netlog.com

Resources

- the great development and it services team at Netlog
- www.netlog.com/go/developer
- www.37signals.com/svn/posts/1509-mr-moore-gets-to-punt-on-sharding
- www.addsimplicity.com/adding_simplicity_an_engi/2008/08/shard-lessons.html
- www.scribd.com/doc/2592098/DVPmysqlucFederation-at-Flickr-Doing-Billions-of-Queries-Per-Day
- startuplessonslearned.blogspot.com/2009/01/sharding-for-startups.html
- www.codefutures.com/weblog/database-sharding
- www.25hoursaday.com/weblog/2009/01/16/BuildingScalableDatabasesProsAndConsOfVariousDatabaseShardingSchemes.aspx
- highscalability.com
- dev.mysql.com/doc/refman/5.1/en/partitioning.html
- www.hibernate.org/414.html
- en.wikipedia.org/wiki/SQLAlchemy

Resources

- spockproxy.sourceforge.net
- www.scribd.com/doc/3865300/Scaling-Web-Sites-by-Sharding-and-Replication
- oracle2mysql.wordpress.com/2007/08/23/scale-out-notes-on-sharding-unique-keys-foreign-keys
- www.flickr.com/photos/kt

Notes to the slides will become available soon at the Netlog Developer Blog and my personal blog (www.jurriaanpersyn.com)